

Bipartite Graph Matching Algorithms for Entity Resolution: An Empirical Evaluation [Experiment, Analysis & Benchmark Papers]

George Papadakis¹, Vasilis Efthymiou², Emanouil Thanos³, Oktie Hassanzadeh⁴

¹National and Kapodistrian University of Athens, Greece gpadadis@di.uoa.gr

²Foundation for Research and Technology - Hellas, Greece vefthym@ics.forth.gr

³Katholieke Universiteit Leuven, Belgium emmanouil.thanos@kuleuven.be

⁴IBM Research AI, USA hassanzadeh@us.ibm.com

ABSTRACT

Entity Resolution (ER) is the task of finding records that refer to the same real-world entities. A common scenario is when entities across two clean sources need to be resolved, which we refer to as Clean-Clean ER. In this paper, we perform an extensive empirical evaluation of 8 bipartite graph matching algorithms that take in as input a bipartite similarity graph and provide as output a set of matched entities. We consider a wide range of matching algorithms, including algorithms that have not previously been applied to ER, or have been evaluated only in other ER settings. We assess the relative performance of the algorithms with respect to accuracy and time efficiency over 10 established, real datasets, from which we extract >700 different similarity graphs. Our results provide insights into the relative performance of these algorithm as well as guidelines for choosing the best one, depending on the ER application at hand.

1 INTRODUCTION

Entity Resolution is a challenging, yet well-studied problem in data integration [5, 6, 23]. A common scenario is Clean-Clean ER (CCER) [41], where the two data sources to be integrated are both clean (i.e., duplicate-free), or are cleaned using single-source entity resolution frameworks. Example applications include Master Data Management [37], where a new clean source needs to be integrated into the clean reference data, and Knowledge Graph matching and completion [19, 48], where an existing clean knowledge base needs to be augmented with an external source.

We focus on methods that take advantage of a large body of work on blocking and matching algorithms, which efficiently compare the entities across two sources and provide as output pairs of entities along with a confidence or similarity score [7, 10, 41]. The output can then be used to decide which pairs should be matched. The simplest approach is specifying as duplicates all the pairs with a score higher than a given threshold. Choosing a single threshold fails to address the issue that in most cases the similarity scores vary significantly depending on the characteristics of the entities. More importantly, for CCER, this approach does not guarantee that each source entity can be matched with at most one other entity. If we view the output as a bipartite *similarity graph*, where the nodes are entity profiles and the edge weights are the matching scores between the candidate duplicates, what we need is finding a *matching* (or independent edge set [30]) so that each entity from one source is matched to at most one entity in the other source.

In this paper, we present the results of our thorough evaluation of efficient bipartite graph matching algorithms for CCER. To the best of our knowledge, our study is the first to primarily focus on bipartite graph matching algorithms, examining the relative

performance of the algorithms in a variety of data sets and methods of creating the input similarity graph. Our goal is to answer the following questions:

- Which bipartite graph matching algorithm is the most accurate one, which is the most robust one, and which offers the best balance between effectiveness and time efficiency?
- How well do the main algorithms scale?
- Which characteristics of the inputs graphs determine the absolute and the relative performance of the algorithms?

By answering these questions we intend to facilitate the selection of the best algorithm for the data at hand.

In summary, we make the following contributions:

- We present an overview of eight efficient bipartite graph matching algorithms along with an analysis of their behavior and complexity. Some of the algorithms are adaptations of efficient graph clustering algorithms that have not been applied to CCER before.
- We organize the input of bipartite graph partitioning algorithms into a taxonomy that is based on the learning-free source of similarity scores/edge weights.
- We assess the relative performance of the matching algorithms with respect to effectiveness and time efficiency.
- We perform an extensive experimental analysis that involves 739 different similarity graphs from 10 established real-world CCER datasets, whose sizes range from several thousands to hundreds of million edges.
- We have publicly released the implementation of all algorithms as well as our experimental results.¹

In what follows, we first discuss the preliminaries and a brief overview of related work in Section 2. We then present the matching algorithms we study in this paper in Section 3. Section 4 describes the generation of similarity graphs. Section 5 presents our experimental evaluation setup and Section 6 presents the results along with a detailed analysis. We conclude the paper with a summary of the key takeaways from our results.

2 PRELIMINARIES

We assume that an *entity profile* or simply *entity* is the description of a real-world object, provided as a set of attribute-value pairs in some *entity collection* V . The problem of Entity Resolution (ER) is to identify such entity profiles (called *matches* or *duplicates*) that correspond to the same real-world object, and place them in a common *cluster* c . In other words, the output of ER, ideally, is a

¹See <https://github.com/scify/JedAIToolkit> for more details.

set of clusters C , each containing all the matching profiles that correspond to a single real-world entity.

In this paper, we focus on the case of *Clean-Clean ER* (CCER), in which we want to match profiles coming from two clean (i.e., duplicate-free) entity collections V_1 and V_2 . This means that the resulting clusters should contain at most two profiles, one from each collection. *Singular clusters*, corresponding to profiles for which no match has been found, are also acceptable.

To generate this clustering, a typical ER pipeline [6] involves the steps of (i) *(meta-)blocking*, i.e., indexing steps that generate *candidate matching pairs*, this way reducing the otherwise quadratic search space of matches, (ii) *matching*, assigning a similarity score to each candidate pair, and (iii) *clustering*, which receives the scored candidate pairs and decides which pairs will be placed together in a cluster. In this work, we evaluate how different clustering methods for CCER perform, when the previous ER steps are fixed.

Problem Definition. The task of **Bipartite Graph Matching** receives as input a bipartite *similarity graph* $SE = (V_1, V_2, E)$, where V_1 and V_2 are two clean entity collections, and $E \subseteq V_1 \times V_2$ is the set of edges with weights in $[0,1]$, corresponding to the similarity scores between entity profiles of the two collections. Its output comprises a set of equivalence clusters C , with each one containing one node $v_i \in V_1 \cup V_2$ or two nodes $v_i \in V_1$ and $v_j \in V_2$ that represent the same real-world object.

Figure 1 (a) shows an example of a bipartite similarity graph, in which node partitions (entity collections) are labeled as A (in orange) and B (in blue). The edges connect only A nodes to B nodes and are associated with a weight that reflects the similarity (matching likelihood) of the connected nodes. Figures 1(b) - (d) show three different outputs of CCER, in which nodes within the same oval (cluster) correspond to matches, i.e., descriptions of the same real-world entity.

2.1 Related work

There is a rich body of literature on various forms of ER [6, 41]. Following the seminal Fellegi-Sunter model for record linkage [12], a major focus of prior work has been on classifying pairs of input records as *match*, *non-match*, or *potential match*. While even some of the early work on record linkage incorporated a 1-1 matching constraint [55], the primary focus of prior work, especially the most recent one, has been on the effectiveness of the classification task, mainly by leveraging Machine [22] and Deep Learning [4, 29, 35].

Inspired by the recent progress and success of prior work on improving the efficiency of ER with blocking and filtering [44], we target ER frameworks where the output of the matching is used to construct a similarity graph that needs to be partitioned for the final step of entity resolution. Hassanzadeh et al. [20] also target such a framework, and perform an evaluation of various graph clustering algorithms for entity resolution. However, they target a scenario where input data sets are not clean or more than two clean sources are merged into a dirty source that contains duplicates in itself, and as a result each cluster could contain more than two records. We refer to this variation of ER as *Dirty ER* [6]. Some of the bipartite matching algorithms we use in this paper are adaptations of the graph clustering algorithms used in [20] for Dirty ER.

Gemmel et al. [15] present two algorithms for CCER as well as more algorithms for different ER settings (e.g., one-to-many and many-to-many). Both algorithms are covered by the clustering algorithms that are included in our study: the MutualFirstChoice is equivalent to our Exact clustering, while the Greedy algorithm is equivalent to UniqueMappingClustering. Finally, the MaxWeight method [15] utilises the exact solution of the maximum weight bipartite matching, for which an efficient heuristic approach is considered in our Best Assignment Heuristic Clustering.

FAMER [48] is a framework that supports multiple matching and clustering algorithms for multi-source ER. Although it studies some common clustering algorithms with those explored in this paper (e.g., Connected Components), the focus on multi-source ER, as opposed to our two-source ER (CCER), makes the direct comparison inapplicable. Additionally, the adaptation of FAMER’s CLIP clustering to work in a CCER setting yields an algorithm equivalent to Unique Mapping Clustering.

Several algorithms have been proposed for weighted bipartite matching in the literature [25, 50, 51]. Our focus in this paper is on efficient algorithms that are significantly faster than the classic Hungarian matching algorithm (also known as the Kuhn-Munkres algorithm[26]), which does not scale to medium-sized datasets of ER, as its time complexity is cubic $O(n^3)$, where n is the total number of nodes in the input graph. An example of more recent matching algorithms that we do not include in this paper due to their time complexity is the work of Schwartz et al. [50] on 1-1 bipartite graph matching with minimum cumulative weights. The authors reduce the problem to a minimum cost flow problem, and use the matching algorithm of Fredman & Tarjan [13] to provide an approximate solution in $O(n^2 \log n)$.

Wang et al. [51] follow a reinforcement learning approach, based on a Q-learning [53] algorithm, for which a state is represented by the pair $(|L|, |R|)$, where $L \subseteq V_1, R \subseteq V_2$ are the nodes matched from the two partitions, and the reward is computed as the sum of the weights of the selected matches. We leave this algorithm outside the scope of this study, as we consider only learning-free methods, but we plan to further explore it in our future works.

Kriege et al. [24] present a linear approximation to the weighted graph matching problem, but for that, they require that the edge weights are assigned by a tree metric, i.e. a similarity measure that satisfies a looser version of the triangle inequality. In this work, we investigate algorithms that are agnostic to such similarity measure properties, assuming only that the weights are in $[0,1]$, as is the case of most existing algorithms.

3 ALGORITHMS

We now describe the main graph matching algorithms for CCER. All of them are configured by a similarity threshold t , which prunes all edges with a lower weight. To describe their time complexity, we denote by $n = |V_1 \cup V_2|$ and $m = |E|$ the number of nodes and edges, respectively, in the bipartite similarity graph $G = (V_1, V_2, E)$ that is given as input. The implementation of all algorithms (in Java) is publicly available through the JedAI toolkit¹ [42].

Connected Components (CNC). This is the simplest algorithm: it discards all edges with a weight lower than the similarity threshold and then computes the transitive closure of the pruned similarity




Figure 1: Example of processing a similarity graph: (a) the similarity graph constructed for a pair of clean entity collections (V_1 in orange and V_2 in blue), (b) the resulting clusters after applying CNC, (c) the resulting partitions/clusters assuming that the approximation algorithms RCA or BAH retrieved the optimal solution for the maximum weight bipartite matching or the assignment problem, and (d) the resulting clusters after applying UMC, BMC or EXC.

graph. In the output, it solely retains the partitions/clusters that contain two entities – one from each entity collection. Using a simple depth-first approach, its time complexity is $O(m)$ [8].

Ricochet Sequential Rippling Clustering (RSR). This algorithm is an adaptation of the homonymous method for Dirty ER in [20] such that it exclusively considers clusters with just one entity from each entity collection. After pruning the edges weighted lower than t , RSR sorts all nodes from both V_1 and V_2 in descending order of the average weight of their adjacent edges. Whenever a new seed is chosen from the sorted list, the first adjacent vertex that is currently unassigned or is closer to the new seed than it is to the seed of its current partition is re-assigned to the new cluster. If a partition is reduced to a singleton after a re-assignment, it is placed in its nearest single-node cluster. The algorithm stops when all nodes have been considered. Its time complexity is $O(nm)$ [54].

Row Column Assignment Clustering (RCA). This approach is based on the Row-Column Scan approximation method in [27] that solves the assignment problem. It requires two passes of the similarity graph, with each pass generating a candidate solution. In the first pass, each entity from V_1 creates a new partition, to which the most similar, currently unassigned entity from V_2 is assigned. Note that, in principle, any pair of entities can be assigned to the same partition at this step even if their similarity is lower than t , since the assignment problem assumes that each vertex from V_1 is connected to all vertices from V_2 (any “job” can be performed by all “men”). The clusters of pairs with similarity less than t are then discarded. In the second pass, the same procedure is applied to the entities/nodes of V_2 . The value of each solution is the sum of the edge weights between the nodes assigned to the same (2-node) partition. The solution with the highest value is returned as output. Its time complexity is $O(|V_1| |V_2|)$ [38].

Best Assignment Heuristic (BAH). This algorithm applies a simple swap-based random-search algorithm to heuristically solve the maximum weight bipartite matching problem and uses the resulting solution to create the output partitions. Initially, each entity from the smaller entity collection is connected to an entity from the larger one. In each iteration of the search process, two entities from the larger entity collection are randomly selected in order to swap their current connections. If the sum of the edge weights of the new pairs is higher than the previous pairs, the swap is accepted.

The algorithm stops when a maximum number of search steps is reached or when a maximum run-time has been exceeded. In our case, the run-time limit has been set to 2 minutes.

Best Match Clustering (BMC). This algorithm is inspired from the Best Match strategy of [33], which solves the stable marriage problem [14], as simplified in BigMat [2]. For each entity of the one entity collection, this algorithm creates a new partition, in which the most similar, not-yet-clustered entity from the other entity collection is also placed – provided that the corresponding edge weight is higher than t . Note that the greedy heuristic for BMC introduced in [33] is the same, in principle, to Unique Mapping Clustering (see below). Note also that BMC is the only algorithm with an additional configuration parameter, apart from the similarity threshold: the entity collection that is used as the basis for creating partitions can be set to V_1 or V_2 . In our experiments, we examine both options and retain the best one. Its time complexity is $O(m)$ [38].

Exact Clustering (EXC). This algorithm is inspired from the Exact strategy of [33]. EXC places two entities in the same partition only if they are mutually the best matches, i.e., the most similar candidates of each other, and their edge weight exceeds t . This approach is basically a stricter, symmetric version of BMC and could also be conceived as a strict version of the reciprocity filter that was employed in [11, 43]. Its time complexity is $O(nm)$.

Király’s Clustering (KRC). This algorithm is an adaptation of the linear time 3/2 approximation to the maximum stable marriage problem, called “New Algorithm” in [21]. Intuitively, the entities of V_1 (“men” [21]) propose to the entities from V_2 with an edge weight higher than t (“women” [21]) to form a partition (“get engaged” [21]). The entities of V_2 accept a proposal under certain conditions (e.g., if it’s the first proposal they receive), and the partitions and preferences are updated accordingly. Entities from V_1 get a second chance to make proposals and the algorithm terminates when all entities of V_1 are in a partition, or no more proposal chances are left. We omit some of the details (e.g., the rare case of “uncertain man”), due to space restrictions, and refer the reader to [21, 38] for more information (e.g., the acceptance criteria for proposals). Its time complexity is $O(n + m \log m)$ [21].

Unique Mapping Clustering (UMC). This algorithm prunes all edges with a weight lower than t , sorts the remaining ones in decreasing weight/similarity and iteratively forms a partition for

the top-weighted pair as long as none of its entities has already been matched to some other. This comes from the *unique mapping constraint* of CCER, i.e., the restriction that each entity from the one entity collection matches at most one entity from the other. Note that the *CLIP Clustering algorithm*, introduced for the multi-source ER problem in [49], is equivalent to UMC in the CCER case that we study. Its time complexity is $O(m \log m)$ [38].

Example. Figure 1 demonstrates an example of applying the above algorithms to the similarity graph in Figure 1(a). For all algorithms, we assume a weight threshold of 0.5.

CNC completely discards the 4-node connected component (A1, B1, A5, B3) and considers exclusively the valid partitions (A2, B2) and (A3, B4), as demonstrated in Figure 1(b).

Algorithms that aim to maximize the total sum of edge weights between the matched entities, such as RCA and BAH, will cluster A1 with B1 and A5 with B3, as shown in Figure 1(c), if they manage to find the optimal solution for the given graph. The reason is that this combination of edge weights yields a sum of $0.6 + 0.6 = 1.2$, which is higher than 0.9, i.e., the sum resulting from clustering A5 with B1 and leaving A1 and B3 as singletons.

UMC starts from the top-weighted edges, matching A5 with B1, A2 with B2 and A3 with B4; A1 and B3 are left as singletons, as shown in Figure 1(d), because their candidates have already been matched to other entities. The same output is produced by EXC, as the entities in each partition consider each other as their most similar candidate. For this reason, BMC also yields the same results assuming that V_2 (blue) is used as the basis entity collection.

The partitions generated by RSR and KRL depend on the sequence of adjacent vertices and proposals, respectively. Given, though, that higher similarities are generally more preferred than increasing total sum by both of these algorithms, the outcome in Figure 1 (d) is the most possible one for these algorithms, too.

4 SIMILARITY GRAPHS

Two types of methods can be used for the generation of the similarity graphs that constitute the input to the above algorithms:

- (i) *learning-free* methods, which produce similarity scores in an unsupervised manner based on the content of the input entities, and
- (ii) *learning-based* methods, which produce probabilistic similarities with the help of a training set [41].

In this work, we exclude the latter, focusing exclusively on learning-free methods. Thus, we make the most of the selected datasets, without sacrificing valuable parts for the construction of the training (and perhaps the validation) set. We also avoid depending on the fine-tuning of numerous configuration parameters, especially in the case of Deep Learning-based methods [52]. Besides, our goal is not to optimize the performance of the CCER process, but to investigate how the main graph matching algorithms perform under a large variety of real settings. For this reason, we produce a large number of similarity graphs per dataset, rather than generating synthetic data.

In this context, we do not apply any blocking method when producing these inputs. Instead, we consider all pairs of entities from different datasets with a similarity higher than 0. This allows for experimenting with a large variety of similarity graph sizes, which range from several thousand to hundreds of million edges.

Besides, the role of blocking, i.e., the pruning of the entity pairs with very low similarity scores, is performed by the similarity threshold t that is employed by all algorithms.

The resulting similarity graphs differ in the number of edges and the corresponding weights, which were produced using different *similarity functions*. Each similarity function consists of two parts: (i) the *representation model*, and (ii) the *similarity measure*.

The **representation model** transforms a textual value into a model that is suitable for applying the selected similarity measure. Depending on the *scope* of these representations, we distinguish them into (i) *schema-agnostic* and (ii) *schema-based*. The former consider all attribute values in an entity description, while the latter consider only the value of a specific attribute. Depending on their *form*, we also distinguish the representations into (i) *syntactic* and (ii) *semantic*. The former operate on the original text of the entities, while the latter operate on vector transformations (embeddings) of the original text that aim to capture the actual connotation, leveraging external information that has been extracted from large and generic corpora through unsupervised learning.

The *schema-based syntactic representations* process each value as a sequence of characters or words and apply to mostly short textual values. For example, the attribute value “Joe Biden” can be represented as the set of tokens {‘Joe’, ‘Biden’}, or the set of character 3-grams {‘Joe’, ‘oe_’, ‘e_B’, ‘_Bi’, ‘Bid’, ‘ide’, ‘den’}.

The *schema-agnostic syntactic representations* process the set of all individual attribute values. We use two types that have been widely applied to document classification tasks [1, 39]:

- (i) an n -gram vector [32], whose dimensions correspond to character or token n -grams and are weighted according to their frequency (TF or TF-IDF score). This approach does not consider the order of n -gram appearances in each value.

- (ii) an n -gram graph [16], which transforms each value into a graph, where the nodes correspond to character or token n -grams, the edges connect those co-occurring in a window of size n and the edge weights denote the n -gram’s co-occurrence frequency. Thus, the order of n -grams in a value is preserved.

Following the previous example, the character 3-gram vector of “Joe Biden” would be a sparse vector with as many dimensions as all the 3-grams appearing in the entity collection and with zeros in all other places except the ones corresponding to the seven character 3-grams of “Joe Biden” listed above. For the places corresponding to those seven 3-grams, the value would be the TF or TF-IDF of each 3-gram. Similarly, a token 2-gram vector of “Joe Biden” would be all zeros, for each token 2-gram appearing in all the values, except for the place corresponding to the 2-gram ‘Joe Biden’, where its value would be 1. A character 3-gram graph would be a graph with seven nodes, one for each 3-gram listed above, connecting the node ‘Joe’ to the nodes ‘oe_’ and ‘e_B’, each with an edge of weight 1, ‘oe_’ to ‘e_B’ and ‘_Bi’, etc. See [38] for more details.

Both approaches build an aggregate representation per entity: the n -gram vectors treat each entity as a “document” and adjust their weights accordingly, while the individual n -gram graphs of each value are merged into a larger “entity graph” through the update operator discussed in [16].

The *semantic representations* treat every text as a sequence of items (words or character n -grams) of arbitrary length and convert

		Scope			
		Schema-agnostic		Schema-based	
		Representation model	Similarity Measure	Representation model	Similarity Measure
Form	Syntactic Similarity	character n-grams (n=2,3,4) and token n-grams (n=1,2,3)	1) Arcs Similarity 2) Cosine Similarity with TF Weights 3) Cosine Similarity with TF-IDF Weights 4) Jaccard Similarity 5) Generalized Jaccard Similarity with TF Weights 6) Generalized Jaccard Similarity with TF-IDF Weights	Character-level	1) Damerau-Levenshtein 2) Levenshtein Distance 3) q-grams Distance 4) Jaro Similarity 5) Needleman Wunch 6)Longest Common Subsequence 7) Longest Common Substring
		character n-gram graphs (n=2,3,4) and token n-gram graphs (n=1,2,3)	1) Containment Similarity 2) Value Similarity 3) Normalized Value Similarity 4) Overall Similarity	Token-level	1) Cosine Similarity 2) Monge-Elkan 3) Block Distance 4) Dice Similarity 5) Overlap Coefficient 6) Euclidean Distance 7) Jaccard Similarity 8) Generalized Jaccard Similarity 9) Euclidean Distance
	Semantic Similarity	fastText and ALBERT	1) Cosine Similarity 2) Euclidean Similarity 3) World Mover's Similarity	fastText and ALBERT	1) Cosine Similarity 2) Euclidean Similarity 3) World Mover's Similarity

Figure 2: Taxonomy of the similarity functions we used to generate the similarity graphs. We use $n \in \{2, 3, 4\}$ for character and $n \in \{1, 2, 3\}$ for token n-grams for both vector and graph models, as in [1, 39]. The graph similarities are defined in [16].

it into a dense numeric vector based on learned external patterns. The closer the connotation of two texts is, the closer are their vectors. These representations come in two main forms, which apply uniformly to schema-agnostic and schema-based settings:

(i) The pre-trained embeddings of word- or character-level. Due to the highly specialized content of ER tasks (e.g., arbitrary alphanumeric in product names), the former, which include *word2vec* [34] and *GloVe* [45], suffer from a high portion of out-of-vocabulary tokens – these are words that cannot be transformed into a vector because they are not included in the training corpora [35]. This drawback is addressed by the character-level embeddings: *fastText* vectorizes a token by summing the embeddings of all its character n-grams [3]. For this reason, we exclusively consider the 300-dimensional *fastText* in the following.

(ii) Transformer-based language models [9] go beyond the shallow, context-agnostic pre-trained embeddings by vectorizing an item based on its context. In this way, they assign different vectors to homonyms, which share the same form, but different meaning (e.g., “bank” as a financial institution or as the border of a river). They also assign similar vectors to synonyms, which have different form, but almost the same meaning (e.g., “enormous” and “vast”). Several BERT-based language models have been applied to ER in [4, 29]. Among them, we exclusively consider the 768-dimensional ALBERT, due to its higher efficiency [28].

Every **similarity measure** receives as input two representation models and produces a score that is proportional to the likelihood that the respective entities correspond to the same real world object: the higher the score, the more similar are the input models and their textual values and the higher is the matching likelihood.

For each type of representation models, we considered a large variety of established similarity measures. We list them in Figure 2, and we formally define them in the Appendix of [38].

5 EXPERIMENTAL SETUP

All experiments were carried out on a server running Ubuntu 18.04.5 LTS with a 32-core Intel Xeon CPU E5-4603 v2 (2.20GHz), 128 GB of RAM and 1.7 TB HDD. All time experiments were executed on a single core. For the implementation of the schema-based syntactic similarity functions, we used the *Simmetrics* Java package². For the schema-agnostic syntactic similarity functions, we used the implementation provided by the *JedAI* toolkit (the implementation of n-gram graphs and the corresponding graph similarities is based on the *JInsect*³ Java toolkit). For the semantic representation models, we employed the Python sister package⁴, which supports both *fastText* and ALBERT. For the computation of the semantic similarities, we used the Python *scipy* package.⁵

Datasets. In our experiments, we use 10 real-world, established datasets for ER, whose technical characteristics appear in Table 1. D_1 , which was introduced in OAEI 2010⁶, contains data about restaurants. D_2 matches products from the online retailers Abt and Buy [23]. D_3 interlinks products from Amazon and the Google Base data API (Google Pr.) [23]. D_4 contains data about publications from DBLP and ACM [23]. $D_5 - D_7$ contain data about television shows from TheTVDB.com (TVDB) and movies from IMDb and themoviedb.org (TMDb) [36]. D_8 contains data about products from Walmart and Amazon [35]. D_9 contains data about scientific publications from DBLP and Google Scholar (Scholar) [23]. D_{10} matches movies from IMDb and DBpedia [40]. All datasets are publicly available through the *JedAI* data repository.⁷

²<https://github.com/Simmetrics/simmetrics>

³<https://github.com/ggianna/JInsect>

⁴<https://pypi.org/project/sister>

⁵<https://www.scipy.org>

⁶<http://oei.ontologymatching.org/2010/im>

⁷<https://github.com/scify/JedAIToolkit/tree/master/data>

Table 1: Technical characteristics of the real datasets for Clean-Clean ER in increasing number of computational cost. $|V_x|$ stands for the number of input entities, $|NVP_x|$ for the total number of name-value pairs, $|A_x|$ for the number of attributes and $|\bar{p}_x|$ for the average number of name-value pairs per entity profile in Dataset_x. $|D(V_1 \cap V_2)|$ denotes the number of duplicates in the ground-truth and $||V_1 \times V_2||$ the number of pairwise comparisons executed by the brute-force approach.

	D ₁	D ₂	D ₃	D ₄	D ₅	D ₆	D ₇	D ₈	D ₉	D ₁₀
Dataset ₁	Rest.1	Abt	Amazon	DBLP	IMDb	IMDb	TMDb	Walmart	DBLP	IMDb
Dataset ₂	Rest.2	Buy	Google Pr.	ACM	TMDb	TVDB	TVDB	Amazon	Scholar	DBpedia
$ V_1 $	339	1,076	1,354	2,616	5,118	5,118	6,056	2,554	2,516	27,615
$ V_2 $	2,256	1,076	3,039	2,294	6,056	7,810	7,810	22,074	61,353	23,182
NVP_1	1,130	2,568	5,302	10,464	21,294	21,294	23,761	14,143	10,064	$1.6 \cdot 10^5$
NVP_2	7,519	2,308	9,110	9,162	23,761	20,902	20,902	$1.14 \cdot 10^5$	$1.98 \cdot 10^5$	$8.2 \cdot 10^5$
$ A_1 $	7	3	4	4	13	13	30	6	4	4
$ A_2 $	7	3	4	4	30	9	9	6	4	7
$ \bar{p}_1 $	3.33	2.39	3.92	4.00	4.16	4.16	3.92	5.54	4.00	5.63
$ \bar{p}_2 $	3.33	2.14	3.00	3.99	3.92	2.68	2.68	5.18	3.24	35.20
$ D(V_1 \cap V_2) $	89	1,076	1,104	2,224	1,968	1,072	1,095	853	2,308	22,863
$ V_1 \times V_2 $	$7.65 \cdot 10^5$	$1.16 \cdot 10^6$	$4.11 \cdot 10^6$	$6.00 \cdot 10^6$	$3.10 \cdot 10^7$	$4.00 \cdot 10^7$	$4.73 \cdot 10^7$	$5.64 \cdot 10^7$	$1.54 \cdot 10^8$	$6.40 \cdot 10^8$

Note that for the schema-based settings (both the syntactic and semantic ones), we used only the attributes that combine high coverage with high distinctiveness. That is, they appear in the majority of entities, while conveying a rich diversity of values, thus yielding high effectiveness. These attributes are for D_1 : “name” and “phone”, for D_2 : “name”, for D_3 : “title”, for D_4 : “title” and “authors”, for D_5 : “modelno” and “title”, for D_6 : “title” and “authors”, for D_7 : “name” and “title”, for D_8 : “title” and “name”, for D_9 : “title” and “abstract”, and for D_{10} : “title”.

Evaluation Measures. In order to assess the relative performance of the above graph matching algorithms, we evaluate both their effectiveness and their time efficiency (and scalability). We measure their *effectiveness*, with respect to a ground truth of known matches, in terms of three measures:

- *Precision* denotes the portion of output partitions that involve two matching entities.
- *Recall* denotes that portion of partitions with two matching entities that are included in the output.
- *F-Measure* (F1) constitutes the harmonic mean of Precision and Recall: $F1 = 2 \cdot \text{Precision} \cdot \text{Recall} / (\text{Precision} + \text{Recall})$.

All measures are defined in $[0, 1]$, with higher values indicating higher effectiveness.

For *time efficiency*, we measure the average run-time of an algorithm for each setting, i.e., the time that intervenes between receiving the weighted similarity graph as input and returning the partitions as output, over 10 repeated executions.

Generation Process. To generate a large variety of input similarity graphs, we apply every similarity function described in Section 4 to all datasets in Table 1. We actually apply all combinations of representation models and similarity measures in Figure 2, thus yielding 60 schema-agnostic syntactic similarity graphs per dataset, 16 schema-based similarity graphs per attribute in each dataset, and 12 semantic similarity graphs per dataset. Note that we did not apply any fine-tuning to ALBERT, as our goal is not optimize ER performance, but rather to produce diverse inputs.

To estimate the algorithms’ performance, we first apply min-max normalization to the edge weights of all similarity graphs,

Table 2: The number of similarity graphs $|G|$ as well as their size, in terms of the average number of edges $|\bar{E}|$, per dataset.

	Syntactic Similarities				Semantic Similarities			
	Schema-based		Schema-ag.		Schema-based		Schema-ag.	
	$ G $	$ \bar{E} \cdot 10^6$	$ G $	$ \bar{E} \cdot 10^6$	$ G $	$ \bar{E} \cdot 10^6$	$ G $	$ \bar{E} \cdot 10^6$
D_1	20	0.16	46	0.72	8	0.26	2	0.76
D_2	12	1.05	47	0.64	2	1.16	2	1.16
D_3	14	2.89	53	2.65	2	4.11	2	4.11
D_4	27	4.49	24	3.84	12	5.99	4	6.00
D_5	24	5.81	48	11.92	12	8.22	2	30.64
D_6	25	8.39	45	10.99	12	12.31	2	39.81
D_7	26	2.80	42	12.21	12	36.44	5	46.99
D_8	26	28.10	47	37.31	2	37.70	-	-
D_9	20	119.18	46	77.56	6	154.26	2	154.36
D_{10}	13	250.73	43	317.17	2	378.51	-	-
Σ	207	-	441	-	70	-	21	-

regardless of the similarity function that produced them, to ensure that they are restricted to $[0, 1]$. Next, we apply every algorithm to every input similarity graph by varying its similarity threshold from 0.05 to 1.0 with a step of 0.05. The largest threshold that achieves the highest F-Measure is selected as the optimal one, determining the performance of the algorithm for the particular input.

Special care was taken to clean the experimental results from noise. We removed all similarity graphs where all matching entities had a zero edge weight. We also removed all noisy graphs, where all algorithms achieve an F-Measure lower than 0.25. Finally, we cleaned our data from duplicate inputs, i.e., similarity graphs that emanate from the same dataset but different similarity function and have the same number of edges, while at least two different algorithms achieve their best performance with the same similarity threshold, exhibiting almost identical effectiveness, i.e., the difference in F-Measure and precision or recall is less than 0.2%.

The characteristics of the retained similarity graphs appear in Table 2. Overall, there are 739 different similarity graphs, most of which rely on syntactic similarity functions and the schema-agnostic settings, in particular. The reason is that the semantic

Table 3: Average performance across all similarity graphs.

	<u>Precision</u>		<u>Recall</u>		<u>F-Measure</u>	
	μ	σ	μ	σ	μ	σ
<i>CNC</i>	0.801	0.185	0.403	0.257	0.490	0.237
<i>RSR</i>	0.615	0.228	0.455	0.239	0.499	0.216
<i>RCA</i>	0.590	0.224	0.502	0.238	0.518	0.211
<i>BAH</i>	0.548	0.236	0.383	0.282	0.408	0.246
<i>BMC</i>	0.631	0.212	0.582	0.221	0.586	0.196
<i>EXC</i>	0.735	0.197	0.544	0.242	0.591	0.199
<i>KRC</i>	0.696	0.200	0.597	0.223	0.619	0.187
<i>UMC</i>	0.645	0.212	0.628	0.212	0.618	0.193

similarities assign relatively high similarity scores to most pairs of entities, thus resulting in poor performance for all considered algorithms – especially in the schema-agnostic settings. Every dataset is represented by at least 58 similarity graphs, in total, while the average number of edges ranges from 160K to 379M. This large set of real-world similarity graphs allows for a rigorous testing of the graph matching algorithms under diverse conditions.

6 EXPERIMENTAL ANALYSIS

In this section, we analyze the experimental results, which we split into effectiveness results and time efficiency results. Our goal is not to declare a single winner, but rather to help identify the use cases for which one algorithm should be preferred over others.

6.1 Effectiveness Measures

The most important performance aspect of clustering algorithms is their ability to effectively distinguish the matching from the non-matching pairs. In this section, we examine this aspect, addressing the following questions:

QE(1): What is the trade-off between precision and recall that is achieved by each algorithm?

QE(2): Which algorithm is the most/least effective?

QE(3): How does the type of input affect the effectiveness of the evaluated algorithms?

QE(4): Which other factors affect their effectiveness?

To answer QE(1) and QE(2), we consider the average performance (μ) of all algorithms across all input similarity graphs, which is reported in Table 3. We observe that all algorithms emphasize on precision at the cost of lower recall. The most balanced algorithm is UMC, as it yields the smallest difference between the two measures (just 0.017). In contrast, CNC constitutes the most imbalanced algorithm, as its precision is almost double its recall. The former achieves the second best F-Measure, being very close to the top performer KRC, while the latter achieves the second worst F-Measure, surpassing only BAH. Note that BAH is the least robust with respect to all measures, as indicated by their standard deviation (σ), due to its stochastic functionality, while CNC, UMC and KRC are the most robust with respect to precision, recall and F-Measure, respectively. Among the other algorithms, EXC and BMC are closer to KRC and UMC, with the former achieving the third highest F1, while RSR and RCA lie closer to CNC, with RSR exhibiting the third lowest F1.

An interesting observation related to QE(2) is that the performance of all algorithms is highly correlated. The Pearson correlation between the precision, recall and F-Measure of all pairs of

algorithms is consistently positive, taking values well above 0.6 in all cases. The correlation between EXC and KRC exceeds 0.96 for all three measures, while UMC is even more correlated with BMC. The only exception is BAH, whose correlation with all other algorithms fluctuates between 0.2 and 0.6, due to its stochastic nature. These patterns indicate that the higher the performance of an algorithm is over a particular similarity graph, the higher will be the performance of any other algorithm over the same input – with the sole exception of BAH.

Another interesting observation is that EXC consistently achieves higher precision and lower recall than BMC. This should be expected, given that EXC requires an additional reciprocity check before declaring that two entities match. We notice, however, that the gain in precision is greater than the loss in recall and, thus, EXC yields a higher F-Measure than BMC, on average. Note also that in the vast majority of cases, the best choice for the entity collection that BMC uses as the basis for creating clusters is the largest one.

To answer QE(3), Figure 3 presents the distribution of precision, recall and F-Measure of all algorithms across the four types of similarity graphs’ origin. For the schema-based syntactic weights, we observe in Figure 3(a) that the average precision of all algorithms increases significantly in comparison to the one in Table 3 - from 4.0% (CNC) to 16.8% (BMC). For CNC and RSR, this is accompanied by an increase in average recall (by 7.7% and 5.1%, respectively), while for all other algorithms, the average recall drops between 4.8% (EXC) and 7.5% (KRC). This means that the schema-based syntactic similarities reinforce the imbalance between precision and recall in Table 3 in favor of the former for all algorithms except CNC and RSR. The average F-Measure drops only for KRC (by 0.2%), which is now outperformed by UMC. Similarly, BMC exceeds EXC in terms of average F1 (0.603 vs 0.599), because the increase in its mean precision is much higher than the decrease in its mean recall. Finally, it is worth noting that this type of input increases significantly the robustness of all algorithms, as the standard deviation of F1 drops by more than 12% for all algorithms, but the stochastic BAH.

The opposite patterns are observed for schema-agnostic syntactic weights in Figure 3(b): the imbalance between precision and recall is reduced, as on average, the former drops from 1.9% (EXC) to 5.3% (BAH), while the latter raises from 2.4% (RSR) to 7.6% (RCA). The imbalance is actually reversed for BMC and UMC, whose average recall (0.613 and 0.664, resp.) exceeds their average precision (0.606 and 0.622, resp.). The only exception is CNC, which increases both its average and average precision. Overall, there are minor, positive changes in the average F-Measure of most algorithms, with KRC and EXC retaining a minor edge over UMC and BMC, respectively.

Regarding the semantic similarity weights, we observe in Figure 3(c) similar patterns with the schema-based syntactic ones: precision increases for all algorithms except CNC and RSR, while recall drops in all cases. In this case, though, the latter change is stronger than the former, leading to lower average F-Measures than those in Table 3. In the case of schema-agnostic semantic weights, all measures in Figure 3(d) drop to a significant extent (>15% in most cases) when compared to Table 3. It is also remarkable that the standard deviation of all measures increases to a significant extent for both schema-based and schema-agnostic weights in relation to their syntactic counterparts, despite the fewer similarity graphs. As a result, the robustness of all algorithms over semantic weights

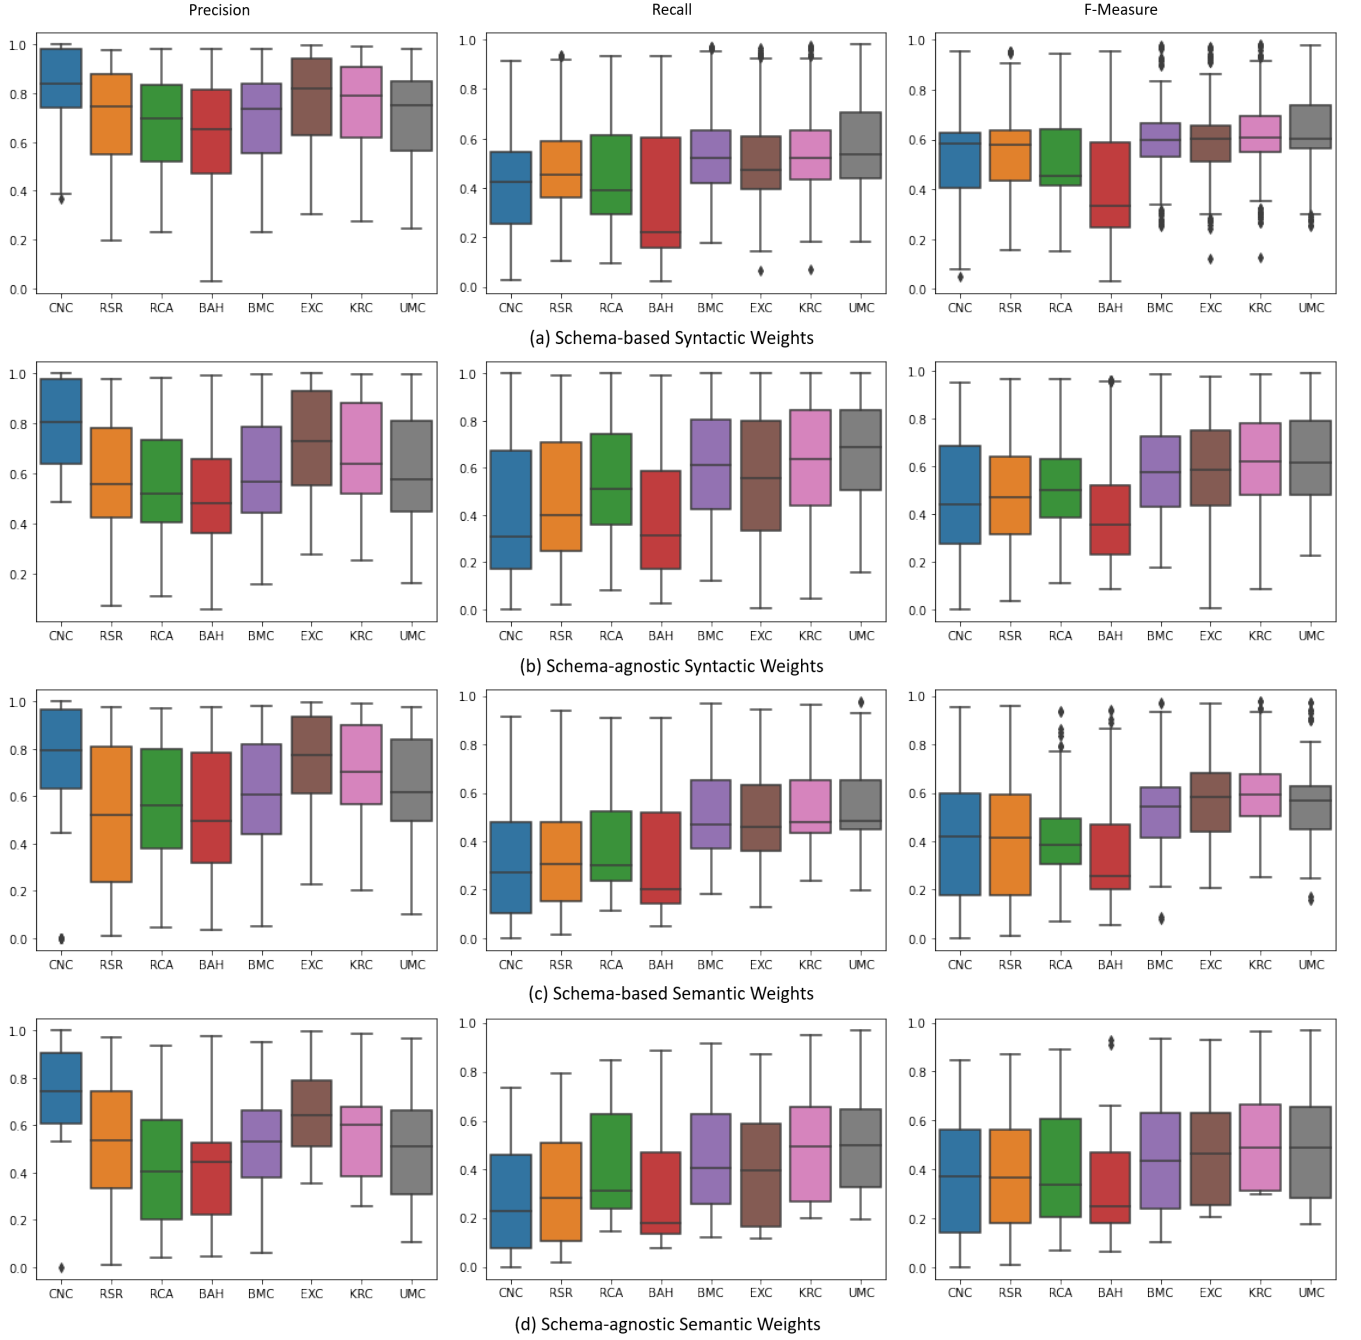


Figure 3: Precision (left), recall (center) and F-Measure (right) of all algorithms over similarity graphs with (a) schema-based syntactic, (b) schema-agnostic syntactic, (c) schema-based semantic, and (d) schema-agnostic semantic edge weights.

is limited. Nevertheless, KRC and EXC maintain a clear lead over UMC and BMC, respectively.

To answer QE(4), we distinguish the similarity graphs into three categories according to the portion of duplicates in their ground truth with respect to the size of $|V_1|$ and $|V_2|$:

(i) *Balanced* (BLC) are the entity collections where the vast majority of entities in V_i are matched with an entity in V_j ($i=1 \wedge j=2$ or

$i=2 \wedge j=1$). This category includes all similarity graphs generated from D_2 , D_4 and D_{10} .

(ii) *One-sided* (OSD) are the entity collections, where only the vast majority of entities in V_1 are matched with an entity from V_2 , or vice versa. OSD includes all graphs stemming from D_3 and D_9 .

(iii) *Scarce* (SCR) are the entity collections, where a small portion of entities in V_i are matched with an entity in V_j ($i=1 \wedge j=2$ or $i=2 \wedge j=1$). This category includes all graphs generated from D_1 , D_5 - D_8 .

Table 4: The number of times each algorithm achieves the highest and second highest F1 for a particular similarity graph, $\#Top1$ and $\#Top2$, resp., as well as the average difference Δ (%) with the second highest F1 across all types of edge weights for balanced (BLC), one-sided (OSD) and scarce (SCR) entity collections. OVL stands for the overall sums or averages across all similarity graphs per category. Note that there are ties for both $\#Top1$ and $\#Top2$: 16 and 40, resp., over schema-based syntactic weights, 17 and 11, resp., over schema-agnostic syntactic weights, 9 and 2, resp., over schema-based semantic weights.

		Syntactic Similarities								Semantic Similarities							
		Schema-based				Schema-agnostic				Schema-based				Schema-agnostic			
		BLC	OSD	SCR	OVL	BLC	OSD	SCR	OVL	BLC	OSD	SCR	OVL	BLC	OSD	SCR	OVL
CNC	$\#Top1$	-	-	18	18	-	-	48	48	-	-	1	1	-	-	-	-
	Δ (%)	-	-	0.41	0.41	-	-	7.59	7.59	-	-	0.33	0.33	-	-	-	-
	$\#Top2$	-	-	8	8	-	-	8	8	-	-	4	4	-	-	-	-
RSR	$\#Top1$	-	-	4	4	-	-	1	1	-	-	1	1	-	-	-	-
	Δ (%)	-	-	1.90	1.90	-	-	0.51	0.51	-	-	0.33	0.33	-	-	-	-
	$\#Top2$	-	-	7	7	-	-	5	5	-	-	1	1	-	-	1	1
RCA	$\#Top1$	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-
	Δ (%)	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-
	$\#Top2$	-	-	-	-	-	-	1	1	-	-	-	-	-	-	-	-
BAH	$\#Top1$	8	-	1	9	40	-	2	42	2	-	1	3	2	-	1	3
	Δ (%)	3.69	-	1.90	3.49	5.55	-	0.46	5.31	12.72	-	0.67	8.70	13.96	-	0.52	9.48
	$\#Top2$	8	-	3	11	7	5	6	18	-	-	2	2	-	-	-	-
BMC	$\#Top1$	-	-	6	6	-	-	7	7	-	-	2	2	-	-	-	-
	Δ (%)	-	-	0.57	0.57	-	-	1.41	1.41	-	-	0.21	0.21	-	-	-	-
	$\#Top2$	2	2	27	31	12	5	18	35	-	-	2	2	-	-	-	-
EXC	$\#Top1$	-	4	35	39	-	11	79	90	-	3	14	17	-	-	5	5
	Δ (%)	-	0.53	0.87	0.84	-	0.18	1.18	1.67	-	1.62	2.54	2.38	-	-	4.18	4.18
	$\#Top2$	-	10	31	41	-	19	72	91	-	3	19	22	-	2	2	4
KRC	$\#Top1$	22	15	43	80	16	47	79	142	11	5	30	46	3	4	4	11
	Δ (%)	1.36	1.60	0.35	0.87	4.15	2.54	4.05	3.56	2.92	1.47	4.92	4.07	4.06	10.32	5.01	6.68
	$\#Top2$	12	17	57	86	39	40	87	166	3	3	11	17	1	-	5	6
UMC	$\#Top1$	22	15	30	67	58	41	29	128	3	-	6	9	1	-	1	2
	Δ (%)	4.99	1.75	1.19	2.56	4.51	3.21	2.51	3.64	2.11	-	0.34	0.93	0.22	-	1.00	0.61
	$\#Top2$	30	5	28	63	56	30	42	128	13	2	9	24	5	2	3	10

We apply this categorization to the four main types of similarity graphs defined by Figure 2 and for each subcategory, we consider three new effectiveness measures:

- (i) $\#Top1$ denotes the number of times an algorithm achieves the maximum F-Measure for a particular category of similarity graphs,
- (ii) Δ (%) stands for the average difference (expressed as a percentage) between the highest and the second highest F1 across all similarity graphs of the same category, and
- (iii) $\#Top2$ denotes the number of times an algorithm scores the second highest F1 for a particular category of similarity graphs.

Note that in case of ties, we increment $\#Top1$ and $\#Top2$ for all involved algorithms. Note also that these three effectiveness measures also allow for answering QE(2) in more detail.

The results for these measures are reported in Table 4. For schema-based syntactic weights, there is a strong competition between KRC and UMC for the highest effectiveness. Both algorithms achieve the maximum F1 for the same number of similarity graphs in the case of balanced and one-sided entity collections. For the former inputs, though, UMC exhibits consistently high performance, as it ranks second in almost all cases that it is not the top performer, unlike KRC, which comes second in 1/3 of these cases. Additionally, UMC achieves significantly higher Δ than KRC. For one-sided entity collections, KRC takes a minor lead over UMC: even though

its Δ is slightly lower, it comes second three times more often than UMC. For scarce entity collections, KRC takes a clear lead over UMC, outperforming it with respect to both $\#Top1$ and $\#Top2$ to a large extent. UMC excels only with respect to Δ .

Among the remaining algorithms, CNC, RSR, BMC and EXC seem suitable only for scarce entity collections. RSR actually achieves the highest Δ , while EXC achieves the second highest $\#Top1$ and $\#Top2$, outperforming UMC. Regarding BAH, we observe that for balanced entity collections, it outperforms all algorithms for 15% of the similarity graphs, achieving the highest Δ and comes second for an equal number of inputs. This is in contrast to the poor average performance reported in Table 3, but is explained by its stochastic nature, which gives rise to an unstable performance, as indicated by the significantly higher σ than all other algorithms for all effectiveness measures.

In the case of schema-agnostic syntactic edge weights, UMC verifies its superiority over KRC for balanced entity collections with respect to all measures. KRC is actually outperformed by BAH, which achieves 2.5 times more often the top F1, while exhibiting the highest Δ among all algorithms. For one-sided entity collections, KRC excels with respect to $\#Top1$ and $\#Top2$, but UMC achieves significantly higher Δ , while EXC constitutes the third best algorithm overall, as for the schema-based syntactic edge weights. In

the case of scarce entity collections, the two competing algorithms are KRC and EXC, as they share the highest $\#Top1$. Yet, the former achieves three times higher Δ and slightly higher $\#Top2$. Surprisingly, CNC ranks second in terms of $\#Top1$, while achieving the highest Δ by far, among all algorithms. As a result, UMC is left at the fourth place, followed by BMC.

For the semantic edge weights, we observe the following patterns: for the balanced entity collections, only KRC, BAH and UMC exhibit the highest performance for both schema-based and schema-agnostic weights. They excel in $\#Top1$, Δ and $\#Top2$, respectively. For one-sided entity collections, KRC is the dominant algorithm, especially in the case of schema-agnostic weights. For the schema-based ones, EXC consistently achieves very high performance, too. For scarce entity collections, there is a strong competition between KRC and EXC; the former consistently outperforms the latter with respect to Δ , while EXC excels in $\#Top1$ for schema-agnostic weights and in $\#Top2$ for schema-based ones.

We examined other patterns with respect to additional characteristics of the entity collections, such as the distribution of positive and negative weights (i.e., between matching and non-matching entities, respectively) and the domain (e-commerce for D_2 , D_3 and D_8 , bibliographic data for D_4 and D_9 as well as movies for D_5 - D_8 and D_{10}). Yet, no clear patterns emerged in these cases.

6.2 Time Efficiency

The (relative) run-time of the evaluated algorithms is a crucial aspect for the task of Entity Resolution, due to the very large similarity graphs, which comprise thousands of entities/nodes and (hundreds of) millions of edges/entity pairs, as reported in Table 2. Below, we study this aspect along with the scalability of the considered algorithms over the 739 different similarity graphs. More specifically, we examine the following questions:

QT(1): Which algorithm is the fastest one?

QT(2): Which factors affect the run-time of the algorithms?

QT(3): How scalable are the algorithms to large input sizes?

The average run-times over 10 executions of the evaluated algorithms per dataset and type of edge weights are reported in Table 5.

Regarding QT(1), we observe that all algorithms are quite fast, as they are all able to process even the largest similarity graphs (i.e., those of D_9 and D_{10}) within minutes or even seconds. CNC is the fastest one almost in all cases, due to the simplicity of its approach. It is followed in close distance by BMC and RSR, with the former consistently outperforming the latter. EXC is also very efficient, but as expected, it is always slower than BMC, due to the additional reciprocity check it involves. On the other extreme lies BAH, which constitutes by far the slowest method, yielding in many cases 2 or even 3 orders of magnitude longer run-times. The reason is the large number (10,000) of search steps we allow per dataset. For the largest datasets, its maximum run-time actually equals the run-time limit of 2 minutes, except for D_9 , where the very large number of entities in V_2 delays the activation of the time-out. The rest of the algorithms lie between these two extremes: KRC is the slowest one, on average, while UMC and RCA exhibit significantly lower run-times. Among the most effective algorithms, EXC is significantly faster than UMC, which is significantly faster than KRC.

Table 5: The average run-time in milliseconds per clustering, dataset and type of edge weights.

	CNC	RSR	RCA	BAH	BMC	EXC	KRC	UMC
D ₁	2	10	7	1,889	2	3	14	4
D ₂	7	22	23	1,149	9	55	77	44
D ₃	17	41	66	2,265	21	29	198	68
D ₄	26	63	157	2,111	34	61	392	273
D ₅	35	87	643	2,734	43	48	344	110
D ₆	51	89	713	2,997	52	56	456	78
D ₇	19	42	932	3,141	24	24	167	27
D ₈	171	257	1,386	122,644	176	203	1,905	476
D ₉	1,277	1,238	2,847	146,630	1,675	20,795	14,925	2,857
D ₁₀	1,391	2,833	24,300	129,113	1,827	2,028	139,458	10,006

(a) Schema-based, syntactic inputs

D ₁	12	19	13	1,888	15	22	51	22
D ₂	7	19	20	1,083	11	17	53	63
D ₃	27	40	58	2,174	37	84	614	506
D ₄	37	55	127	2,163	39	40	362	84
D ₅	80	238	756	3,429	155	273	989	1,007
D ₆	72	137	843	3,731	89	184	651	179
D ₇	82	141	1,075	3,748	98	138	615	290
D ₈	289	347	719	124,063	315	408	2,729	300
D ₉	437	740	2,038	146,301	528	1,613	5,920	2,556
D ₁₀	1,861	8,429	24,845	129,942	3,394	19,457	190,731	159,643

(b) Schema-agnostic, syntactic inputs

D ₁	9	16	44	1,819	17	20	42	27
D ₂	8	24	16	1,029	10	30	122	71
D ₃	28	46	124	2,210	225	472	2,312	3,566
D ₄	35	171	82	2,244	156	255	1,250	1,369
D ₅	76	456	1462	2,874	230	375	1,295	1,637
D ₆	118	752	628	3,202	188	684	2,135	2,781
D ₇	26	126	550	2,872	91	145	521	669
D ₈	236	286	9,242	122,843	258	3,382	2,036	368
D ₉	833	15,311	17,140	140,556	2,262	20,164	41,960	49,823
D ₁₀	2,650	4,506	3,291	127,030	11,824	50,030	340,066	131,259

(c) Schema-based, semantic inputs

D ₁	14	64	108	1,640	15	16	76	24
D ₂	8	35	26	959	48	91	275	473
D ₃	23	36	207	1,836	29	336	3,256	4,010
D ₄	34	62	152	1,884	46	54	2,733	218
D ₅	171	242	883	2,202	202	504	1,746	406
D ₆	436	305	1,496	2,545	565	1,010	3,716	3,684
D ₇	207	865	1,510	2,972	2,384	3,432	8,156	13,743
D ₈	-	-	-	-	-	-	-	-
D ₉	824	1,138	402,886	143,788	1,140	23,784	51,590	4,339
D ₁₀	-	-	-	-	-	-	-	-

(d) Schema-agnostic, semantic inputs

Regarding QT(2), there are two main factors that affect the reported run-times: (i) the time complexity of the algorithms, and (ii) the similarity threshold used for pruning the search space. Regarding the first factor, we observe that the run-times in Table 5 verify the time complexities described in Section 3. With $O(m)$, CNC and BMC are the fastest ones, followed by RSR and EXC with $O(nm)$, RCA with $O(|V_1| |V_2|)$, UMC with $O(m \log m)$ and KRC




Figure 4: Scalability analysis of all algorithms over all similarity graphs with (i) schema-based syntactic, (ii) schema-agnostic syntactic, (iii) schema-based semantic and (iv) schema-agnostic semantic edge weights. The horizontal axis corresponds to the number of edges in the similarity graphs and the vertical one to the run-time in milliseconds (maximum value=16.7 min).

with $O(n + m \log m)$. BAH’s run-time is determined by the number of search steps and the run-time limit.

More interesting is the effect of the similarity thresholds: the higher their optimal value (i.e., the one maximizing F1) is, the fewer edges are retained in the similarity graph and the faster is its

processing. The optimal similarity threshold depends on various aspects, such as the clustering algorithm and the type of edge weights. This means that the relative time efficiency of algorithms with the same theoretical complexity should be attributed to their different similarity threshold. For example, the average optimal thresholds

for CNC and BMC over all schema-based syntactic weights are 0.755 and 0.669, respectively, while over schema-agnostic syntactic weights they are 0.409 and 0.327, respectively. These large differences account for the consistently lower run-time of CNC in both cases. The larger the difference in the similarity threshold, the larger is the difference in the run-times. The same applies to RSR and EXC, as their optimal thresholds amount to 0.758 and 0.627, respectively, for schema-based syntactic weights and to 0.406 and 0.291, respectively, for their schema-agnostic counterparts.

Note that the similarity threshold typically accounts for the relative run-times that conflict with the relative time complexities, too: in case an algorithm runs faster than another one with lower time complexity, this is typically caused by the higher similarity threshold it employs. For example, KRC exhibits a lower average run-time than EXC over D_9 for schema-based syntactic weights, because their mean optimal similarity thresholds amount to 0.550 and 0.490, respectively. Similarly, UMC runs much faster than EXC over D_8 with schema-agnostic syntactic weights, because their mean optimal similarity thresholds amount to 0.427 and 0.387, respectively.

Finally, the similarity threshold accounts for the relative run-times between the same algorithm over different types of edge weights. For example, EXC is 5 times slower over the schema-agnostic syntactic weights of D_{10} than their schema-based counterparts, even though the former involve just 25% more edges than the latter, as reported in Table 2. This difference should be attributed to the huge difference in the average optimal thresholds: 0.153 for the former weights and 0.535 for the latter ones. The same applies to UMC, whose average run-time increases by 16 times when comparing the schema-based with the schema-agnostic weights, because its average optimal threshold drops from 0.481 to 0.110.

To answer QT(3), Figure 4 presents the scalability analysis of every algorithm over all similarity graphs for each type of edge weights. In each diagram, every point corresponds to the run-time of a different similarity graph. We observe that for all algorithms, the run-time increases linearly with the size of the similarity graphs: as the number of edges increases by four orders of magnitude, from 10^4 to 10^8 , the run-times increase to a similar extent in most cases. For all algorithms, though, there are outlier points that deviate from the “central” curve. The larger the number of outliers is, the less robust is the time efficiency of the corresponding algorithm, due to its sensitivity to the size of the graph and the similarity threshold. In this respect, the least robust algorithms are RSR over schema-based syntactic weights and UMC with EXC over schema-agnostic syntactic weights. These patterns seem to apply to the semantic weights, too, despite the limited number of the similarity graphs, especially in the case of schema-agnostic weights.

Note that there are two exceptions to these patterns, namely RCA and BAH. The diagrams of the former algorithm seem to involve a much lower number of points, as its time complexity depends exclusively on the number of entities in the input entity collections, i.e., $O(|V_1| + |V_2|)$. As a result, different similarity graphs from the same dataset yield similar run-times that coincide in the diagrams of Figure 4. Regarding BAH, it exhibits a step-resembling scalability graph, because its processing terminates after a pre-defined timeout or a fixed number of iterations (whichever comes first), independently of the size of the similarity graph.

7 CONCLUSIONS

The most important patterns we draw from our experiments are the following:

(i) There is no clear winner among the considered algorithms, as the best performing one for a particular similarity graph depends, among others, on the type of edge weights and the portion of duplicates with respect to the total number of nodes/entities.

(ii) CNC constitutes the fastest algorithm, due to its simplicity and the high similarity thresholds it employs, achieving the highest precision at the cost of low recall. It frequently outperforms all other algorithms with respect to F1 in the case of scarce entity collections with syntactic weights, especially the schema-agnostic ones.

(iii) RSR is a fast algorithm that rarely achieves high effectiveness - in the case of scarce entity collections.

(iv) RCA is an efficient method that never excels in effectiveness.

(v) BAH constitutes a slow, stochastic approach that is capable of the best and the worst. It frequently achieves the highest by far F1 over balanced entity collections (and rarely over scarce ones), but in most cases, it yields the lowest scores with respect to all effectiveness measures.

(vi) BMC is the second fastest algorithm that tries to balance precision and recall, being particularly effective in the case of scarce entity collections - especially in combination with syntactic weights.

(vii) EXC improves BMC by boosting precision at the cost of lower recall and higher run-time. It consistently achieves (close to) the maximum F1 over scarce and one-sided entity collections, losing only to KRC and (rarely) to UMC. Given, though, that it outperforms both algorithms to a significant extent with respect to run-time, it constitutes the best choice for applications requiring both high effectiveness and efficiency/scalability.

(viii) KRC achieves very high or the highest effectiveness in most cases, especially over one-sided and scarce entity collections. This comes, though, at the cost of higher (yet stable) run-times than its top-performing counterparts.

(ix) UMC is the best choice for balanced entity collections, especially when coupled with syntactic weights, exhibiting a much more robust performance than BAH. It achieves very high (and frequently the highest) effectiveness in the rest of the cases, too. Its run-time, though, is rather unstable, depending largely on the optimal similarity threshold.

Our focus in this work has been on entity resolution in an unsupervised setting where large enough training data is not available. Following the approach of Reas et al. [46], an interesting avenue for future work is a supervised learning system incorporating a variety of bipartite graph matching algorithms in a unified framework and effectively learning the right parameters for a given data set.

REFERENCES

- [1] Fotis Aisopos, George Papadakis, Konstantinos Tserpes, and Theodora A. Varvarigou. 2012. Content vs. context for sentiment analysis: a comparative analysis over microblogs. In *ACM Conference on Hypertext and Social Media*. 187–196.
- [2] Ali Assi, Hamid Mcheick, and Wajdi Dhifli. 2019. BIGMAT: A Distributed Affinity-Preserving Random Walk Strategy for Instance Matching on Knowledge Graphs. In *IEEE Big Data*. 1028–1033.
- [3] Piotr Bojanowski, Edouard Grave, Armand Joulin, and Tomáš Mikolov. 2017. Enriching Word Vectors with Subword Information. *Trans. Assoc. Comput. Linguistics* 5 (2017), 135–146.
- [4] Ursin Brunner and Kurt Stockinger. 2020. Entity Matching with Transformer Architectures - A Step Forward in Data Integration. In *EDBT*. 463–473.

- [5] Peter Christen. 2012. *Data Matching - Concepts and Techniques for Record Linkage, Entity Resolution, and Duplicate Detection*. Springer.
- [6] Vassilis Christophides, Vasilis Efthymiou, Themis Palpanas, George Papadakis, and Kostas Stefanidis. 2021. An Overview of End-to-End Entity Resolution for Big Data. *ACM Comput. Surv.* 53, 6 (2021), 127:1–127:42.
- [7] Vassilis Christophides, Vasilis Efthymiou, and Kostas Stefanidis. 2015. *Entity Resolution in the Web of Data*. Morgan & Claypool Publishers.
- [8] Sanjoy Dasgupta, Christos H. Papadimitriou, and Umesh V. Vazirani. 2008. *Algorithms*. McGraw-Hill.
- [9] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In *NAACL-HLT*. 4171–4186.
- [10] Xin Luna Dong and Divesh Srivastava. 2015. *Big Data Integration*. Morgan & Claypool Publishers.
- [11] Vasilis Efthymiou, George Papadakis, Kostas Stefanidis, and Vassilis Christophides. 2019. MinoanER: Schema-Agnostic, Non-Iterative, Massively Parallel Resolution of Web Entities. In *EDBT*. OpenProceedings.org, 373–384.
- [12] I. P. Fellegi and A. B. Sunter. 1969. A Theory for Record Linkage. *J. Amer. Statist. Assoc.* 64, 328 (1969), 1183–1210.
- [13] Michael L. Fredman and Robert Endre Tarjan. 1987. Fibonacci heaps and their uses in improved network optimization algorithms. *J. ACM* 34, 3 (1987), 596–615.
- [14] D. Gale and L. S. Shapley. 1962. College Admissions and the Stability of Marriage. *Am. Math. Mon.* 69, 1 (1962), 9–15.
- [15] Jim Gemmell, Benjamin I. P. Rubinstein, and Ashok K. Chandra. 2011. Improving Entity Resolution with Global Constraints. *CoRR* abs/1108.6016 (2011).
- [16] George Giannakopoulos, Vangelis Karkaletsis, George A. Vouras, and Panagiotis Stamatopoulos. 2008. Summarization system evaluation revisited: N-gram graphs. *ACM Trans. Speech Lang. Process.* 5, 3 (2008), 5:1–5:39.
- [17] George Giannakopoulos and Themis Palpanas. 2010. Content and type as orthogonal modeling features: a study on user interest awareness in entity subscription services. *International Journal of Advances on Networks and Services* 3, 2 (2010).
- [18] Osamu Gotoh. 1982. An improved algorithm for matching biological sequences. *Journal of molecular biology* 162, 3 (1982), 705–708.
- [19] Claudio Gutierrez and Juan F. Sequeda. 2020. *Knowledge Graphs: A Tutorial on the History of Knowledge Graph's Main Ideas*. Association for Computing Machinery, 3509–3510. <https://doi.org/10.1145/3340531.3412176>
- [20] Oktie Hassanzadeh, Fei Chiang, Renée J. Miller, and Hyun Chul Lee. 2009. Framework for Evaluating Clustering Algorithms in Duplicate Detection. *Proc. VLDB Endow.* 2, 1 (2009), 1282–1293.
- [21] Zoltán Király. 2013. Linear Time Local Approximation Algorithm for Maximum Stable Marriage. *Algorithms* 6, 3 (2013), 471–484.
- [22] Pradap Konda, Sanjib Das, Paul Suganthan G. C., AnHai Doan, Adel Ardalan, Jeffrey R. Ballard, Han Li, Fatemah Panahi, Haojun Zhang, Jeffrey F. Naughton, Shishir Prasad, Ganesh Krishnan, Rohit Deep, and Vijay Raghavendra. 2016. Magellan: Toward Building Entity Matching Management Systems. *Proc. VLDB Endow.* 9, 12 (2016), 1197–1208.
- [23] Hanna Köpcke, Andreas Thor, and Erhard Rahm. 2010. Evaluation of entity resolution approaches on real-world match problems. *Proc. VLDB Endow.* 3, 1 (2010), 484–493.
- [24] Nils M. Kriege, Pierre-Louis Giscard, Franka Bause, and Richard C. Wilson. 2019. Computing Optimal Assignments in Linear Time for Approximate Graph Matching. In *ICDM*. 349–358.
- [25] Harold W. Kuhn. 2010. The Hungarian Method for the Assignment Problem. In *50 Years of Integer Programming 1958-2008 - From the Early Years to the State-of-the-Art*. Springer, 29–47.
- [26] H. W. Kuhn and Bryn Yaw. 1955. The Hungarian method for the assignment problem. *Naval Res. Logist. Quart.* (1955), 83–97.
- [27] Jerome M Kurtzberg. 1962. On approximation methods for the assignment problem. *Journal of the ACM (JACM)* 9, 4 (1962), 419–439.
- [28] Zhenzhong Lan, Mingda Chen, Sebastian Goodman, Kevin Gimpel, Piyush Sharma, and Radu Soricut. 2020. ALBERT: A Lite BERT for Self-supervised Learning of Language Representations. In *ICLR*.
- [29] Yuliang Li, Jinfeng Li, Yoshihiko Suhara, Jin Wang, Wataru Hirota, and Wang-Chiew Tan. 2021. Deep Entity Matching: Challenges and Opportunities. *ACM J. Data Inf. Qual.* 13, 1 (2021), 1:1–1:17.
- [30] L. Lovasz and M. D. Plummer. [n.d.]. *Matching theory*.
- [31] Fragkiskos D. Malliaros, Polykarpos Meladianos, and Michalis Vazirgiannis. 2018. Graph-based Text Representations: Boosting Text Mining, NLP and Information Retrieval with Graphs. In *WWW Tutorials*.
- [32] Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze. 2008. *Introduction to information retrieval*. Cambridge University Press.
- [33] Sergey Melnik, Hector Garcia-Molina, and Erhard Rahm. 2002. Similarity Flooding: A Versatile Graph Matching Algorithm and Its Application to Schema Matching. In *ICDE*. 117–128.
- [34] Tomás Mikolov, Ilya Sutskever, Kai Chen, Gregory S. Corrado, and Jeffrey Dean. 2013. Distributed Representations of Words and Phrases and their Compositionality. In *NIPS*. 3111–3119.
- [35] Sidharth Mudgal, Han Li, Theodoros Rekatsinas, AnHai Doan, Youngchoon Park, Ganesh Krishnan, Rohit Deep, Esteban Arcaute, and Vijay Raghavendra. 2018. Deep Learning for Entity Matching: A Design Space Exploration. In *SIGMOD*. 19–34.
- [36] Daniel Obaczka, Jonathan Schuchart, and Erhard Rahm. 2021. EA-GER: Embedding-Assisted Entity Resolution for Knowledge Graphs. *CoRR* abs/2101.06126 (2021).
- [37] Boris Otto and Andreas Reichert. 2010. Organizing Master Data Management: Findings from an Expert Survey. In *Proceedings of the 2010 ACM Symposium on Applied Computing (SAC)*. 106–110. <https://doi.org/10.1145/1774088.1774111>
- [38] George Papadakis, Vasilis Efthymiou, Emanouil Thanos, and Oktie Hassanzadeh. 2021. Bipartite Graph Matching Algorithms for Entity Resolution: An Empirical Evaluation (Extended Version). <https://github.com/scify/JedAIToolkit/blob/master/documentation/bipartiteGraphMatchingExtendedVersion.pdf>.
- [39] George Papadakis, George Giannakopoulos, and Georgios Paliouras. 2016. Graph vs. bag representation models for the topic classification of web documents. *World Wide Web* 19, 5 (2016), 887–920.
- [40] George Papadakis, Ekaterini Ioannou, Claudia Niederée, and Peter Fankhauser. 2011. Efficient entity resolution for large heterogeneous information spaces. In *WSDM*. 535–544.
- [41] George Papadakis, Ekaterini Ioannou, Emanouil Thanos, and Themis Palpanas. 2021. *The Four Generations of Entity Resolution*. Morgan & Claypool Publishers.
- [42] George Papadakis, Georgios M. Mandilaras, Luca Gagliardelli, Giovanni Simonini, Emmanouil Thanos, George Giannakopoulos, Sonia Bergamaschi, Themis Palpanas, and Manolis Koubarakis. 2020. Three-dimensional Entity Resolution with JedAI. *Inf. Syst.* 93 (2020), 101565.
- [43] George Papadakis, George Papastefanatos, Themis Palpanas, and Manolis Koubarakis. 2016. Boosting the Efficiency of Large-Scale Entity Resolution with Enhanced Meta-Blocking. *Big Data Res.* 6 (2016), 43–63.
- [44] George Papadakis, Dimitrios Skoutas, Emmanouil Thanos, and Themis Palpanas. 2020. Blocking and Filtering Techniques for Entity Resolution: A Survey. *ACM Comput. Surv.* 53, 2 (2020), 31:1–31:42. <https://doi.org/10.1145/3377455>
- [45] Jeffrey Pennington, Richard Socher, and Christopher D. Manning. 2014. Glove: Global Vectors for Word Representation. In *EMNLP*. 1532–1543.
- [46] Russell Reas, Steve Ash, Rob Barton, and Andrew Borthwick. 2018. SuperPart: Supervised Graph Partitioning for Record Linkage. In *IEEE International Conference on Data Mining, ICDM 2018, Singapore, November 17-20, 2018*. IEEE Computer Society, 387–396. <https://doi.org/10.1109/ICDM.2018.00054>
- [47] François Rousseau and Michalis Vazirgiannis. 2013. Graph-of-word and TW-IDF: new approach to ad hoc IR. In *CIKM*. 59–68.
- [48] Alieh Saeedi, Markus Nentwig, Eric Peukert, and Erhard Rahm. 2018. Scalable Matching and Clustering of Entities with FAMER. *Complex Syst. Informatics Model. Q.* 16 (2018), 61–83.
- [49] Alieh Saeedi, Eric Peukert, and Erhard Rahm. 2018. Using Link Features for Entity Clustering in Knowledge Graphs. In *ESWC (Lecture Notes in Computer Science)*, Vol. 10843. Springer, 576–592.
- [50] Justus Schwartz, Angelika Steger, and Andreas Weißl. 2005. Fast Algorithms for Weighted Bipartite Matching. In *WEA (Lecture Notes in Computer Science)*, Vol. 3503. 476–487.
- [51] Yansheng Wang, Yongxin Tong, Cheng Long, Pan Xu, Ke Xu, and Weifeng Lv. 2019. Adaptive Dynamic Bipartite Graph Matching: A Reinforcement Learning Approach. In *ICDE*. 1478–1489.
- [52] Zhengyang Wang, Bunyamin Sisman, Hao Wei, Xin Luna Dong, and Shuiwang Ji. 2020. CorDEL: A Contrastive Deep Learning Approach for Entity Linkage. In *ICDM*.
- [53] Christopher J. C. H. Watkins and Peter Dayan. 1992. Technical Note Q-Learning. *Mach. Learn.* 8 (1992), 279–292.
- [54] Derry Tanti Wijaya and Stéphane Bressan. 2009. Ricochet: A family of unconstrained algorithms for graph clustering. In *International Conference on Database Systems for Advanced Applications*. Springer, 153–167.
- [55] W. E. Winkler. 2006. *Overview of Record Linkage and Current Research Directions*. Technical Report. Bureau of the Census.

Algorithm 1: Ricochet SR Clustering (RSR)

Input: Similarity Graph $G = (V_1, V_2, E)$, similarity threshold t
Output: A set of Partitions $C = \{c_1, c_2, \dots, c_n\}$

```

1  $C \leftarrow \emptyset$ 
2  $Center \leftarrow \emptyset$ 
3 foreach  $v \in (V_1 \cup V_2)$  do // Initialization
4    $simWithCenter(v) \leftarrow 0$ 
5    $Partition(v) \leftarrow \emptyset$ 
6    $centerOf(v) \leftarrow v$ 
7  $Q \leftarrow G.nodesInDecWeight(v, w(v))$ 
   //  $w(v) = \sum_{e \in adj(v)} e.sim / |adj(v)|$ 
8 while  $Q \neq \emptyset$  do
9    $v_i \leftarrow Q.pop()$  // the vertex with highest weight
10   $ToReassign \leftarrow \emptyset$ 
11  foreach  $e = (v_i, v_j, sim) \in E : sim > t$  do // for  $v_i$ 's
    adjacent edges
12    if  $v_j \in Center$  then
13      continue
14    if  $e.sim > simWithCenter(v_j)$  then
15       $Partition(centerOf(v_j)).remove(v_j)$  // remove
         $v_j$  from its previous partition
16       $Partition(v_i) \leftarrow Partition(v_i) \cup v_j$ 
17       $ToReassign \leftarrow ToReassign \cup centerOf(v_j)$  // it
        is now a singleton
18       $simWithCenter(v_j) \leftarrow e.sim$ 
19       $centerOf(v_j) \leftarrow v_i$ 
20      break
21  if  $|Partition(v_i)| > 0$  then
22    if  $centerOf(v_i) \neq v_i$  then // if  $v_i$  was previously
        in another partition
23       $Partition(centerOf(v_i)).remove(v_i)$ 
24       $ToReassign \leftarrow ToReassign \cup centerOf(v_i)$ 
25       $Center \leftarrow Center \cup v_i$ 
26       $Partition(v_i) \leftarrow Partition(v_i) \cup v_i$  // put  $v_i$  in its
        partition
27       $centerOf(v_i) \leftarrow v_i$ 
28       $simWithCenter(v_i) \leftarrow 1$ 
29  foreach  $v_k \in ToReassign$  do
30     $maxSim \leftarrow 0$ 
31     $cMax \leftarrow v_k$ 
32    foreach  $e = (v_k, v_\ell, sim) \in E : sim > t$  do // find
        singleton with the highest similarity with  $v_k$ 
        to reassign it
33      if  $e.sim > maxSim$  and  $|Partition(v_\ell)| < 2$  then
34         $cMax \leftarrow v_\ell$ 
35         $maxSim \leftarrow e.sim$ 
36      if  $maxSim > 0$  then
37         $Partition(v_k) \leftarrow \emptyset$ 
38         $Partition(cMax) \leftarrow Partition(cMax) \cup v_k$ 
39  foreach  $v_i \in (V_1 \cup V_2)$  do
40    if  $|Partition(v_i)| > 0$  then
41       $C \leftarrow C \cup Partition(v_i)$ 
42  return  $C$ 

```

Algorithm 2: Connected Components (CNC)

Input: Similarity Graph $G = (V_1, V_2, E)$, similarity threshold t
Output: A set of clusters $C = \{c_1, c_2, \dots, c_n\}$

```

1 foreach  $e = (v_i, v_j, sim) \in E$  do
2   if  $sim < t$  then
3      $E \leftarrow E - (e)$ 
4  $C_1 \leftarrow transitiveClosure(G)$ 
5  $C \leftarrow \emptyset$ 
6 foreach  $c_i \in C_1$  do
7   if  $|c_i| == 2 \wedge c_i \cap V_1 \neq \emptyset \wedge c_i \cap V_2 \neq \emptyset$  then
8      $C \leftarrow C \cup \{c_i\}$ 
9 return  $C$ 

```

APPENDIX**A ALGORITHMS**

Ricochet Sequential Rippling Clustering (RSR). This algorithm, outlined in Algorithm 1, is an adaptation of the homonymous method for Dirty ER in [20] such that it exclusively considers partitions with just one entity from each input dataset. Initially, RSR sorts all nodes from both input datasets in descending order of the average weight of their adjacent edges (Line 7). Whenever a new seed is chosen from the sorted list, we consider all its adjacent edges with a weight higher than t (Lines 8-11). The first adjacent vertex that is currently unassigned or is closer to the new seed than it is to the seed of its current partition is re-assigned to the new partition (Lines 14-16). If a partition is reduced to a singleton after a re-assignment, either because the chosen vertex (Line 17) or the seed (Line 24) was previously in it, it is placed in its nearest single-node partition (Lines 30-39).

The algorithm stops when all nodes have been considered. In the worst case the algorithm has to iterate through n vertices and each time reassign n vertices to their most similar adjacent vertex, therefore its time complexity is $O(n \cdot m)$ [54].

Connected Components (CNC). This is the simplest algorithm for bipartite graph matching and is outlined in Algorithm 2. First, it discards all edges with a weight lower than the given similarity threshold (Lines 1-3). Then, it computes the transitive closure of the pruned similarity graph (Line 4). In the output, it solely retains the partitions that contain two entities – one from each input dataset (Lines 6-8). Using a simple depth-first approach, its time complexity is $O(n + m) \sim O(m)$, given that $m \gg n$ [8].

Row Column Assignment Clustering (RCA). This approach, outlined in Algorithm 3, is based on the Row-Column Scan approximation method in [27] that solves the assignment problem. It requires two passes of the similarity graph, with each pass generating a candidate solution. In the first pass, each entity from the source dataset creates a new partition, to which the most similar, currently unassigned entity from the target dataset is assigned (Lines 7-17). Note that, in principle, any pair of entities can be assigned to the same partition at this step even if their similarity is lower than t , since the assignment problem assumes that each vertex from V_1 is connected to all vertices from V_2 (any "job" can be performed

Algorithm 3: Row Column Clustering (RCA)

Input: Similarity Graph $G = (V_1, V_2, E)$, similarity threshold t
Output: A set of partitions $C = \{c_1, c_2, \dots, c_n\}$

```

1  $C_1 \leftarrow \emptyset$ 
2  $C_2 \leftarrow \emptyset$ 
3  $M_1 \leftarrow \emptyset$  // matched nodes from  $V_1$ 
4  $M_2 \leftarrow \emptyset$  // matched nodes from  $V_2$ 
5  $D_1 \leftarrow 0$  // assignment value of  $C_1$ 
6  $D_2 \leftarrow 0$  // assignment value of  $C_2$ 
7 foreach  $v_i \in V_1$  do
8    $c_i \leftarrow \{v_i\}$  // create a new partition containing  $v_i$ 
9    $Q_i \leftarrow V_2(sim(v_i))$  // a priority queue of  $V_2$ 's nodes
    in decreasing sim with  $v_i$ 
10  while  $Q_i \neq \emptyset$  do
11     $v_2 \leftarrow Q_i.pop()$ 
12    if  $v_2 \notin M_2$  then // if  $v_2$  is not yet matched
13       $c_i \leftarrow c_i \cup \{v_2\}$  // add  $v_2$  to partition  $c_i$ 
14       $M_2 \leftarrow M_2 \cup \{v_2\}$ 
15       $D_1 \leftarrow D_1 + \{sim(v_1, v_2)\}$ 
16      break
17   $C_1 \leftarrow C_1 \cup \{c_i\}$ 
18 foreach  $v_j \in V_2$  do
19    $c_j \leftarrow \{v_j\}$  // create a new partition containing  $v_j$ 
20    $Q_j \leftarrow V_1(sim(v_j))$  // a priority queue of  $V_1$ 's nodes
    in decreasing sim with  $v_j$ 
21   while  $Q_j \neq \emptyset$  do
22      $v_1 \leftarrow Q_j.pop()$ 
23     if  $v_1 \notin M_1$  then // if  $v_1$  is not yet matched
24        $c_j \leftarrow c_j \cup \{v_1\}$  // add  $v_1$  to partition  $c_j$ 
25        $M_1 \leftarrow M_1 \cup \{v_1\}$ 
26        $D_2 \leftarrow D_2 + \{sim(v_1, v_2)\}$ 
27       break
28    $C_2 \leftarrow C_2 \cup \{c_j\}$ 
29 if  $D_1 > D_2$  then // get maximal assignment
30    $C = C_1$ 
31 else
32    $C = C_2$ 
33 foreach  $c = \{v_i, v_j\} \in C$  do
34   if  $sim(v_i, v_j) < t$  then // check similarities
35      $C = C \setminus c$  // remove partition pairs with
    similarity less than  $t$ 
36 return  $C$ 

```

by all "men"). In the second pass, the same procedure is applied to the entities/nodes of the target dataset (Lines 18-28). The value of each solution is the sum of the edge weights between the nodes assigned to the same (2-node) partition (Lines 17,28). The solution with the highest value is returned as output, after discarding the pairs with similarity less than t (Lines 29-36).

At each pass the algorithm iterates over all nodes/entities of one of the entity collection from the other entity with maximum similarity from the other entity collection. Therefore, its time complexity is $O(|V_1| |V_2|)$.

Best Assignment Heuristic (BAH). This algorithm applies a simple swap-based random-search algorithm to heuristically solve the

Maximum Weight Bipartite Matching problem and uses the resulting solution to create the output partitions. Its functionality is outlined in Algorithm 4. Initially, each entity from the smaller input dataset is connected to an entity from the larger input dataset (Line 9). In each iteration of the search process (Line 10), two entities from the larger dataset are randomly selected (Lines 12-13) in order to swap their current connections. If the sum of the edge weights of the new pairs is higher than the previous pairs (Line 15-19), the swap is accepted (Lines 20-24). The algorithm stops when a maximum number of search steps is reached or when a maximum run-time has been exceeded. In our case, the run-time has been set to 2 minutes.

Algorithm 4: Best Assignment Heuristic (BAH)

Input: Similarity Graph $G = (V_1, V_2, E) : |V_1| > |V_2|$, similarity threshold t , max number of moves $maxNumMoves$
Output: A set of partitions $C = \{c_1, c_2, \dots, c_n\}$

```

1  $C \leftarrow \emptyset$ 
2  $numMoves \leftarrow 0$ 
3 foreach  $(v_i^1, v_j^2) \in (V_1 \times V_2)$  do
4    $d(v_i^1, v_j^2) \leftarrow 0$  // initialize pair contributions
5 foreach  $e = (v_i^1, v_j^2, sim) \in E$ , with  $e.sim > t$  do
6    $d(v_i^1, v_j^2) \leftarrow sim$  // initialize pair contributions
7 foreach  $v_i^1 \in V_1, v_i^2 \in V_2 : i \leq |V_2|$  do
8    $c_i \leftarrow \{v_i^1, v_i^2\}$  // initialize partitions
9    $p(v_i^1) = v_i^2$ 
10 while  $numMoves < maxNumMoves$  do
11    $numMoves \leftarrow numMoves + 1$ 
12    $i = nextRand(|V_1|)$ 
13    $j = nextRand(|V_1|) : j \neq i$ 
14    $D \leftarrow 0$ 
15   if  $p(v_i^1) \neq null$  then // check swaps
16      $D \leftarrow d(v_j^1, p(v_i^1)) - d(v_i^1, p(v_i^1))$ 
17   if  $p(v_j^1) \neq null$  then // check swaps
18      $D \leftarrow D + d(v_i^1, p(v_j^1)) - d(v_j^1, p(v_j^1))$ 
19   if  $D \geq 0$  then // if swaps increase assignment value
20      $temp \leftarrow p(v_j^1)$  // perform swaps
21      $p(v_j^1) \leftarrow p(v_i^1)$ 
22      $p(v_i^1) \leftarrow temp$ 
23      $c_i \leftarrow \{v_i^1, p(v_i^1)\}$ 
24      $c_j \leftarrow \{v_j^1, p(v_j^1)\}$ 
25 return  $C$ 

```

Best Match Clustering (BMC). This algorithm is inspired from the Best Match strategy of [33], which solves the Stable Marriage problem [14], as simplified in BigMat [2]. Its functionality is outlined in Algorithm 5. For each entity of the one dataset, this algorithm creates a new partition (Lines 4-5), in which the most similar, not-yet-clustered entity from the other dataset is also placed - provided that the corresponding edge weight is higher than t (Lines 6-12). Note that the greedy heuristic for BMC introduced in [33] is the same, in principle, to Unique Mapping Clustering (see below). Note also that BMC is the only algorithm with an additional configuration parameter, apart from the similarity threshold: the input

dataset that is used as the basis for creating partitions can be set to the source or the target dataset. In our experiments, we examine both options and retain the best one.

The algorithm iterates over the nodes of one of the datasets searching for its adjacent vertex with maximum similarity, therefore its time complexity is $O(m)$.

Algorithm 5: Best Match Clustering (BMC)

Input: Similarity Graph $G = (V_1, V_2, E)$, similarity threshold t
Output: A set of clusters $C = \{c_1, c_2, \dots, c_n\}$

```

1  $C \leftarrow \emptyset$ 
2  $M_2 \leftarrow \emptyset$  // matched nodes from  $V_2$ 
3 foreach  $v_i \in V_1$  do
4    $c_i \leftarrow \{v_i\}$  // create a new cluster containing  $v_i$ 
5    $Q_i \leftarrow v_i.edgesDecOrder(t)$  // edges in desc. sim  $> t$ 
6   while  $Q_i \neq \emptyset$  do
7      $e \leftarrow Q_i.pop()$ 
8     if  $e.v_2 \notin M_2$  then // if  $v_2$  is not yet matched
9        $c_i \leftarrow c_i \cup \{e.v_2\}$  // add  $v_2$  to cluster  $c_i$ 
10       $M_2 \leftarrow M_2 \cup \{e.v_2\}$ 
11      break
12    $C \leftarrow C \cup \{c_i\}$ 
13 return  $C$ 

```

Exact Clustering (EXC). This algorithm is inspired from the Exact strategy of [33]. Its functionality is outlined in Algorithm 6. Initially, it creates an empty priority queue for every vertex (Lines 3-6). Then, it populates the queue of every vertex v_i with all its adjacent edges that exceed the given similarity threshold t , sorting them in decreasing weight (Lines 7-9). Subsequently, EXC places two entities in the same partition (Lines 10-18) only if they are mutually the best matches, i.e., the most similar candidates of each other (Line 15). This approach is basically a stricter, symmetric version of BMC and could also be conceived as a strict version of the reciprocity filter that was employed in [11, 43].

Its time complexity is $O(nm)$, since the algorithm iterates over each vertex of one dataset searching for its adjacent vertex with maximum similarity and then performs the same search for the latter vertex.

Unique Mapping Clustering (UMC). This algorithm is outlined in Algorithm 8. Initially, it iterates over all edges and those with a weight higher than t are placed in a priority queue that sorts them in decreasing weight/similarity (Lines 5-7). Subsequently, it iteratively forms a partition (Line 11) for the top-weighted pair (Line 9), as long as none of its entities has already been matched to some other (Line 10). This approach relies on the *unique mapping constraint* of CCER, i.e., the restriction that each entity from the one input dataset matches at most one entity from the other. Note that the *CLIP Clustering algorithm*, introduced for the multi-source ER problem in [49], is equivalent to UMC when there are only two input datasets (i.e., in the CCER case that we study).

Its time complexity is $O(m \log m)$, due to the cost that is required for sorting all edges.

Király's Clustering (KRC). This algorithm is an adaptation of the linear time 3/2 approximation to the Maximum Stable Marriage problem, called "New Algorithm" in [21]. Intuitively, the entities of

the source dataset ("men" [21]) propose to the entities (Line 16) from the target dataset with an edge weight higher than t ("women" [21]) to form a partition ("get engaged" [21]). Its functionality is outlined in Algorithm 7. The entities of the target dataset accept a proposal under certain conditions (e.g., if it's the first proposal they receive - Line 17), and the partitions and preferences are updated accordingly (Lines 18, 22, 24). Entities from the source dataset get a second chance to make proposals (Lines 5, 27-30) and the algorithm terminates when all entities of the first dataset are in a partition (Line 13), or no more proposal chances are left (Line 27). We omit some of the details (e.g., the rare case of "uncertain man"), due to space restrictions, and refer the reader to [21] for more information (e.g., the acceptance criteria for proposals). Its time complexity is $O(n + m \log m)$ [21].

Algorithm 6: Exact Clustering (EXC)

Input: Similarity Graph $G = (V_1, V_2, E)$, similarity threshold t
Output: A set of clusters $C = \{c_1, c_2, \dots, c_n\}$

```

1  $C \leftarrow \emptyset$ 
2  $M_2 \leftarrow \emptyset$  // matched nodes from  $V_2$ 
3 foreach  $v_i \in V_1$  do
4    $Q1_i \leftarrow \emptyset$  // initialize a PQ in desc. sim
5 foreach  $v_j \in V_2$  do
6    $Q2_j \leftarrow \emptyset$  // initialize a PQ in desc. sim
7 foreach  $e = (v_i, v_j, sim) \in E$ , with  $e.sim > t$  do
8    $Q1_i.push(e)$ 
9    $Q2_j.push(e)$ 
10 foreach  $v_i \in V_1$  do
11    $c_i \leftarrow \{v_i\}$  // create a new cluster containing  $v_i$ 
12    $e \leftarrow Q1_i.pop()$  // the best edge for  $v_i$ 
13    $v_j \leftarrow e.v_j$  // the best match for  $v_i$  is  $v_j$ 
14    $e_2 \leftarrow Q2_j.pop()$  // the best edge for  $v_j$ 
15   if  $e_2.v_i = v_i$  then // if the best match for  $v_j$  is  $v_i$ 
16      $c_i \leftarrow c_i \cup \{v_j\}$  // add  $v_j$  to cluster  $c_i$ 
17      $M_2 \leftarrow M_2 \cup \{e.v_j\}$ 
18      $C \leftarrow C \cup \{c_i\}$ 
19 return  $C$ 

```

B SIMILARITY FUNCTIONS

In this section, we provide more details about the similarity measures mentioned in Figure 2.

- (1) For the schema-based syntactic representations, which involve short textual values, we considered 16 established similarity measures: Cosine Similarity, Block Distance, Levenshtein Distance, Damerau-Levenshtein Distance, Euclidean Distance, Jaccard Similarity, Generalized Jaccard Similarity, Dice Similarity, Overlap Coefficient, Jaro Similarity, Longest Common Subsequence, Longest Common Substring, Monge-Elkan Similarity, Needleman-Wunch, q-grams Distance, and Simon White Similarity.
- (2) For the n-gram vectors, we used six similarity measures: Cosine and Generalized Jaccard Similarity with both TF and TF-IDF weights, Enhanced Jaccard Similarity with TF weights and ARCS Similarity with TF-IDF weights.

Algorithm 8: Unique Mapping Clustering

Input: Similarity Graph $G = (V_1, V_2, E)$, similarity threshold t
Output: A set of clusters $C = \{c_1, c_2, \dots, c_n\}$

```

1  $C \leftarrow \emptyset$ 
2  $M_1 \leftarrow \emptyset$  // matched nodes from  $V_1$ 
3  $M_2 \leftarrow \emptyset$  // matched nodes from  $V_2$ 
4  $Q \leftarrow \emptyset$ 
5 foreach  $e = (v_i, v_j, sim) \in E$  do
6   if  $e.sim > t$  then
7      $Q.put(e)$  // a PQ of edges in desc. sim >  $t$ 
8 while  $Q \neq \emptyset$  do
9    $e \leftarrow Q.pop()$  // the entity pair with highest sim
10  if  $e.v_i \notin M_1$  and  $e.v_j \notin M_2$  then //  $v_i, v_j$  not matched
11     $C \leftarrow C \cup \{e.v_i, e.v_j\}$ 
12     $M_1 \leftarrow M_1 \cup \{e.v_i\}$ 
13     $M_2 \leftarrow M_2 \cup \{e.v_j\}$ 
14 return  $C$ 

```

Algorithm 7: Király's Clustering

Input: Similarity Graph $G = (V_1, V_2, E)$, similarity threshold t
Output: A set of clusters $C = \{c_1, c_2, \dots, c_n\}$

```

1  $C \leftarrow \emptyset$ 
2  $freeM \leftarrow \emptyset$  // an initially empty linked list
3 foreach  $v_i \in V_1$  do //  $V_1$  corresponds to men in [21]
4    $Q1_i \leftarrow \emptyset$  //  $v_i$ 's edges in desc. sim >  $t$ 
5    $lastChance[i] \leftarrow false$ 
6    $freeM.addLast(v_i)$  // keeps insertion order
7 foreach  $v_j \in V_2$  do //  $V_2$  corresponds to women in [21]
8    $Q2_j \leftarrow \emptyset$  //  $v_j$ 's edges in desc. sim >  $t$ 
9    $fiancé[j] \leftarrow null$ 
10 foreach  $e = (v_i, v_j, sim) \in E$ , with  $e.sim > t$  do
11    $Q1_i.push(e)$ 
12    $Q2_j.push(e)$ 
13 while  $freeM \neq \emptyset$  do
14    $v_i \leftarrow freeM.removeFirst()$  // in insertion order
15   if  $Q1_i \neq \emptyset$  then
16      $v_j \leftarrow Q1_i.pop()$  //  $v_i$ 's preference is  $v_j$ 
17     if  $fiancé[j] = null$  then //  $v_j$  is free
18        $C \leftarrow C \cup \{v_i, v_j\}$  // match  $v_i$  to  $v_j$ 
19     else
20        $v'_i \leftarrow fiancé[j]$  //  $v_j$  was engaged to  $v'_i$ 
21       if  $acceptsProposal(v_j, v_i)$  then // refer to [21]
22          $C \leftarrow C \setminus \{v'_i, v_j\}$  //  $v'_i$  and  $v_j$  break up
23          $freeM.addLast(v'_i)$  //  $v'_i$  is free again
24          $C \leftarrow C \cup \{v_i, v_j\}$  // match  $v_i$  to  $v_j$ 
25          $fiancé[j] \leftarrow v_i$  //  $v_j$  gets engaged to  $v_i$ 
26   else
27     if  $lastChance(v_i) = false$  then
28        $lastChance[i] \leftarrow true$  // 2nd chance for  $v_i$ 
29        $Q1_i \leftarrow recoverInitialQueue(v_i)$ 
30        $freeM.addLast(v_i)$ 
31 return  $C$ 

```

(3) For the n-gram graphs, we used four graph similarity measures: Containment, Value, Normalized Value and Overall Similarity (i.e., the average of the three measures) [16].

(4) For the semantic models, we consider Cosine Similarity, Euclidean Similarity ($=1/(1+Euclidean\ distance)$) and World Mover's Similarity ($=1/(1+World\ Mover's\ Distance)$).

Each category of similarity functions is described in more detail in the following.

B.1 Schema-based syntactic functions

For this category, we use the following similarity and distance measures, as defined and implemented in Simmetrics.

B.1.1 Character-level measures. The following similarity measures are applied to two strings s_1 and s_2 at character level.

Levenshtein Distance: Counts the (minimum) number of insert, delete and substitute operations required to transform one string into the other.

Damerau-Levenshtein Distance: Damerau-Levenshtein Distance only differs to Levenshtein Distance by including transpositions among the operations allowed.

Jaro Similarity: The Jaro Similarity of two strings s_1 and s_2 is given by the formula:

$$similarity(s_1, s_2) = \begin{cases} 0 & , \text{ if } m = 0 \\ \frac{1}{3} \left(\frac{m}{|s_1|} + \frac{m}{|s_2|} + \frac{m-t}{m} \right) & , \text{ else,} \end{cases}$$

where m is the number of common characters, and t is the number of transpositions.

Needleman-Wunch: This similarity measure is the result of applying an algorithm that assigns three scores (seen as parameters) to two sequences of characters s_1 and s_2 , depending on whether aligned characters are a match, a mismatch, or a gap. A match occurs when the two aligned characters are the same, a mismatch when they are not the same, and a gap when for the aligned we need an insert or delete operation. The match, mismatch, gap scores used in this study, as in Simmetrics, are 0, -1, and -2, respectively.

q-grams Distance: It applies a Block Distance (see below) similarity metric over all tri-grams in a string.

Longest Common Substring Similarity: As the name suggests, this measure counts the size of the longest common substring (lcs_{str}) between two strings, divided by the size of the longest string: $similarity(s_1, s_2) = |lcs_{str}(s_1, s_2)| / \max(|s_1|, |s_2|)$.

Longest Common Subsequence Similarity: The difference between this measure and the previous is that a subsequence does not need to consist of consecutive characters: $similarity(s_1, s_2) = |lcs_{seq}(s_1, s_2)| / \max(|s_1|, |s_2|)$.

B.1.2 Word-level measures. The following similarity measures are applied to two strings a and b that are treated as sets or multisets (bags) of words.

Cosine Similarity: The similarity is defined as the cosine of the angle between the multisets (bags) of words a and b expressed as sparse vectors. $similarity(a, b) = a \cdot b / (||a|| ||b||)$.

Euclidean Distance: Compares the frequency of occurrence of each word w in two strings a and b $distance(a, b) = ||a - b|| = \sqrt{\sum_w (freqA(w) - freqB(w))^2}$

Block Distance: Also known as *L1 Distance*, *City Block Distance* and *Manhattan Distance* between two multisets (bags) of words a

and b is the sum of the absolute differences of the frequency of each word in a vs in b : $distance(a, b) = ||a - b||_1$.

Overlap Coefficient: The size of the intersection divided by the smaller of the size of the two sets of words: $similarity(a, b) = |a \cap b| / \min(|a|, |b|)$.

Dice Similarity: The Dice Similarity is defined as twice the shared information (intersection) divided by sum of cardinalities of the two sets of words: $similarity(a, b) = 2|a \cap b| / (|a| + |b|)$.

Simon White Similarity: This similarity is the same as Dice Similarity, with the only difference being that it considers a and b as multisets (bags) of words.

Jaccard Similarity: Computes the size of the intersection divided by the size of the union for two sets of words $similarity(a, b) = |a \cap b| / |a \cup b|$.

Generalized Jaccard Similarity: Same as the Jaccard Similarity, except that the Generalized Jaccard Similarity considers multisets (bags) of words, instead of sets.

Monge-Elkan Similarity: This similarity is the average similarity of the most similar words between two sets of words a and b : $similarity(a, b) = \frac{1}{|a|} \sum_{w_i \in a} \max_{w_j \in b} (sim(w_i, w_j))$, where sim is the optimized Smith-Waterman algorithm [18] that operates as the secondary character-level similarity to compute the similarity of individual words.

B.2 Schema-agnostic syntactic functions

For this category, we use the following similarity functions, as defined and implemented in JedAI.

B.2.1 Bag Models [32]. There are two types of n-grams, the character and the token ones. These give rise to two types of bag models: the *character n-grams model* and the *token n-grams model*. Collectively, they are called *bag* or *vector space models*, because they model an entity e_i as a vector with one dimension for every distinct n-gram in an entity collection E : $BM(e_i) = (w_{i1}, \dots, w_{im})$, where m stands for the *dimensionality* of E (i.e., the number of distinct n-grams in it), while w_{ij} is the weight of the j^{th} dimension that quantifies the importance of the corresponding n-gram for e_i .

The most common weighting schemes are:

(i) **Term Frequency (TF)** sets weights in proportion to the number of times the corresponding n-grams appear in the values of entity e_i . More formally, $TF(t_j, e_i) = f_j / N_{e_i}$, where f_j stands for the occurrence frequency of t_j in e_i , while N_{e_i} is the number of n-grams in e_i , normalizing TF so as to mitigate the effect of different lengths on the weights.

(ii) **Term Frequency-Inverse Document Frequency (TF-IDF)** discounts the TF weight for the most common tokens in the entire entity collection E , as they typically correspond to noise (i.e., stop words). Formally, $TF-IDF(t_j, e_i) = TF(t_j, e_i) \cdot IDF(t_j)$, where $IDF(t_j)$ is the inverse document frequency of the n-gram t_j , i.e., $IDF(t_j) = \log |E| / (|\{e_k \in E : t_j \in e_k\}| + 1)$. In this way, high weights are given to n-grams with high frequency in e_i , but low frequency in E .

To construct the bag model for a specific entity, we aggregate the vectors corresponding to each one of its attribute values. The end result is a weighted vector $(a_i(w_1), \dots, a_i(w_m))$, where $a_i(w_j)$

is the sum of weights, i.e., $a_i(w_j) = \sum_{a_k \in A_i} w_{kj}$, where a_k stands for an individual attribute value in the set of values A_i of entity e_i .

To compare two bag models, $BM(e_i)$ and $BM(e_j)$, one of the following similarity measures is typically used:

(i) **ARCS Similarity (ARCS)** sums the inverse Document Frequency of the common n-grams in two bag models. That is, the rarer the common n-grams are, the higher gets the overall similarity. Formally: $ARCS(BM(e_i), BM(e_j)) = \sum_{k \in BM(e_i) \cap BM(e_j)} \log 2 / \log(DF_1(k) \cdot DF_2(k))$.

(ii) **Cosine Similarity (CS)** measures the cosine of the angle of the weighted vectors. Formally, it is equal to their dot product similarity, normalized by the product of their magnitudes:

$$CS(BM(e_i), BM(e_j)) = \sum_{k=1}^m w_{ik} w_{jk} / (||BM(e_i)|| \cdot ||BM(e_j)||)$$

(iii) **Jaccard Similarity (JS)** treats the document vectors as sets, with weights higher than (equal to) 0 indicating the presence (absence) of the corresponding n-gram. On this basis, it defines as similarity the ratio between the sizes of set intersection and union: $JS(BM(e_i), BM(e_j)) = |BM(e_i) \cap BM(e_j)| / |BM(e_i) \cup BM(e_j)|$.

(iv) **Generalized Jaccard Similarity (GJS)** extends JS so that it takes into account the weights associated with every n-gram:

$$GJS(BM(e_i), BM(e_j)) = \sum_{k=1}^m \min(w_{ik}, w_{jk}) / \sum_{k=1}^m \max(w_{ik}, w_{jk})$$

Both CS and GJS apply seamlessly to both TF and TF-IDF weights.

B.2.2 Graph Models [16, 47]. Recent works suggest that graph models outperform the bag ones in various tasks [31], from Information Retrieval [47] to Document Classification [39]. There are two graph models, one for each type of n-grams, i.e., *token n-gram graphs* [47] and *character n-gram graphs* [16]. Both models represent each entity e_i as an undirected graph G_i that contains one vertex for each n-gram in the attribute values of e_i . An edge connects every pair of vertices/n-grams that co-occur within a window of size n in the values of e_i . Every edge is weighted according to the co-occurrence frequency of the corresponding n-grams. Thus, the graphs incorporate *contextual information* in the form of n-grams' closeness.

To construct the model for an entity, we merge the graphs of its attribute values using the update operator, which is described in [16, 17]. To compare graph models, we can use the following graph similarity measures [16]:

(i) **Containment Similarity (CoS)** estimates the number of edges shared by two graph models, G_i and G_j , regardless of the corresponding weights (i.e., it merely estimates the portion of common n-grams in the original texts). Formally:

$$CoS(G_i, G_j) = \sum_{e \in G_i} \mu(e, G_j) / \min(|G_i|, |G_j|), \text{ where } |G| \text{ is the size of graph } G, \text{ and } \mu(e, G) = 1 \text{ if } e \in G, \text{ or } 0 \text{ otherwise.}$$

(ii) **Value Similarity (VS)** extends CoS by considering the weights of common edges. Formally, using w_e^k for the weight of edge e in G_k : $VS(G_i, G_j) = \sum_{e \in (G_i \cap G_j)} \frac{\min(w_e^i, w_e^j)}{\max(w_e^i, w_e^j) \cdot \max(|G_i|, |G_j|)}$.

(iii) **Normalized Value Similarity (NS)** extends VS by mitigating the impact of imbalanced graphs, i.e., the cases where the comparison between a large graph with a much smaller one yields similarities close to 0. Formally:

$$NS(G_i, G_j) = \sum_{e \in (G_i \cap G_j)} \min(w_e^i, w_e^j) / \max(w_e^i, w_e^j) / \min(|G_i|, |G_j|)$$

(iv) **Overall Similarity (OS)** constitutes the average of the above graph similarity measures, which are all defined $[0, 1]$. Formally: $OS(G_i, G_j) = (CoS(G_i, G_j) + VS(G_i, G_j) + NS(G_i, G_j)) / 3$.